

# Assignment 2

```
In [8]: import numpy as np
import pickle
import matplotlib.pyplot as plt
import time
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.model_selection import GridSearchCV

# Load the emnist_train.pkl file
with open('emnist_train.pkl', 'rb') as f:
    train_dict = pickle.load(f, encoding='bytes')

# Extract the data and labels arrays from the dictionary
train_data = train_dict['data']
train_labels = train_dict['labels']

# Load the emnist_test.pkl file
with open('emnist_test.pkl', 'rb') as f:
    test_dict = pickle.load(f, encoding='bytes')

# Extract the data and labels arrays from the dictionary
test_data = test_dict['data']
test_labels = test_dict['labels']

# Convert the data and labels arrays to NumPy arrays
train_data = np.array(train_data, dtype=np.float32)
train_labels = np.array(train_labels, dtype=np.int32)
test_data = np.array(test_data, dtype=np.float32)
test_labels = np.array(test_labels, dtype=np.int32)

# Define class names
# Create class names for all 62 classes
class_names = ['Class {}'.format(i) for i in range(62)]
#Plot a grid of EMNIST examples of a specified size."""

def plot_examples(data, pred, target, n_rows=5, n_cols=5):
    """Plot a grid of EMNIST examples of a specified size."""
    # Size figure depending on the size of the grid
    plt.figure(figsize=(n_cols * 1.2, n_rows * 2.0))

    for row in range(n_rows):
        for col in range(n_cols):
            # Get the next index of the image
            index = n_cols * row + col

            # Reshape the flattened image to its original shape
            image = data[index].reshape((28, 28))

            # Plot the image at the appropriate place in the grid
            plt.subplot(n_rows, n_cols, index + 1)
            plt.imshow(image, cmap="binary")
            plt.axis('off')
            plt.title(class_names[target[index]] + '\n' + '>>' + class_names[pred[index]])
```


```

plt.tight_layout()
plt.show()

# Plot examples from different classes
fig, axes = plt.subplots(8, 8, figsize=(12, 12))
for i, ax in enumerate(axes.flat):
    if i < len(class_names):
        class_index = i # Compute the class index
        # Get samples of the current class
        class_samples = train_data[train_labels == class_index]
        # Randomly select a sample from the class
        sample_index = np.random.randint(class_samples.shape[0])
        ax.imshow(class_samples[sample_index].reshape((28, 28)), cmap='gray')
        ax.axis('off')
        # Set the title as the class name
        ax.set_title(class_names[class_index])
    else:
        ax.axis('off')

plt.tight_layout()
plt.show()

```



MLP without Hyperparameter tuning

```
In [9]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.utils import to_categorical
import time
import numpy as np
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, precision_score, recall_score

# Reshape the data to 2D for MLP
train_data = train_data.reshape((-1, 28*28))
test_data = test_data.reshape((-1, 28*28))

# Normalize the data
train_data = train_data / 255.0
test_data = test_data / 255.0

# Convert the labels to one-hot encoding
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

# Create the MLP model
model = Sequential([
    Dense(128, activation='relu', input_shape=(28*28,)),
    Dense(64, activation='relu'),
    Dense(len(train_labels[0]), activation='softmax') # Number of output classes
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Record the start time
start_time = time.time()

# Train the model
model.fit(train_data, train_labels, epochs=10, batch_size=128, validation_split=0.1)

# Record the end time
end_time = time.time()

# Calculate and print the training time
train_time = end_time - start_time
print('Training time:', train_time)

# Evaluate the model on test data
test_loss, test_acc = model.evaluate(test_data, test_labels)
print('Test accuracy:', test_acc)

# Generate predictions
preds = model.predict(test_data)
preds_classes = np.argmax(preds, axis=1)

# Convert back to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate and print the classification report
classification_report_results = classification_report(test_labels_int, preds_classes)
print("Classification Report:")
print(classification_report_results)

# Generate and print the confusion matrix
confusion_matrix_results = confusion_matrix(test_labels_int, preds_classes)
print("Confusion Matrix:")
print(confusion_matrix_results)
```

```
accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro')
recall = recall_score(test_labels_int, preds_classes, average='macro')

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
```

Epoch 1/10  
704/704 [=====] - 3s 3ms/step - loss: 1.3098 - accuracy: 0.6502 - val\_loss: 0.9197 - val\_accuracy: 0.7303  
Epoch 2/10  
704/704 [=====] - 2s 3ms/step - loss: 0.7866 - accuracy: 0.7592 - val\_loss: 0.7544 - val\_accuracy: 0.7701  
Epoch 3/10  
704/704 [=====] - 2s 3ms/step - loss: 0.6760 - accuracy: 0.7859 - val\_loss: 0.6960 - val\_accuracy: 0.7813  
Epoch 4/10  
704/704 [=====] - 2s 3ms/step - loss: 0.6141 - accuracy: 0.8003 - val\_loss: 0.6556 - val\_accuracy: 0.7929  
Epoch 5/10  
704/704 [=====] - 2s 3ms/step - loss: 0.5753 - accuracy: 0.8100 - val\_loss: 0.6289 - val\_accuracy: 0.8002  
Epoch 6/10  
704/704 [=====] - 3s 4ms/step - loss: 0.5444 - accuracy: 0.8176 - val\_loss: 0.6218 - val\_accuracy: 0.7970  
Epoch 7/10  
704/704 [=====] - 3s 4ms/step - loss: 0.5191 - accuracy: 0.8245 - val\_loss: 0.6010 - val\_accuracy: 0.8048  
Epoch 8/10  
704/704 [=====] - 2s 3ms/step - loss: 0.5027 - accuracy: 0.8281 - val\_loss: 0.6022 - val\_accuracy: 0.8059  
Epoch 9/10  
704/704 [=====] - 2s 3ms/step - loss: 0.4850 - accuracy: 0.8336 - val\_loss: 0.5930 - val\_accuracy: 0.8042  
Epoch 10/10  
704/704 [=====] - 2s 3ms/step - loss: 0.4680 - accuracy: 0.8367 - val\_loss: 0.5933 - val\_accuracy: 0.8078  
Training time: 23.027072191238403  
625/625 [=====] - 1s 2ms/step - loss: 0.5916 - accuracy: 0.8069  
Test accuracy: 0.8069499731063843  
625/625 [=====] - 1s 788us/step

Classification Report:

	precision	recall	f1-score	support
0	0.64	0.81	0.72	976
1	0.67	0.90	0.77	1023
2	0.89	0.96	0.92	1003
3	0.94	0.97	0.96	1035
4	0.90	0.95	0.92	903
5	0.80	0.92	0.85	928
6	0.92	0.96	0.94	959
7	0.97	0.97	0.97	1098
8	0.93	0.93	0.93	941
9	0.94	0.89	0.91	929
10	0.84	0.89	0.87	170
11	0.70	0.79	0.74	118
12	0.73	0.77	0.75	316
13	0.79	0.69	0.74	128
14	0.90	0.80	0.85	162
15	0.84	0.80	0.82	261
16	0.71	0.53	0.61	60
17	0.81	0.85	0.83	74
18	0.62	0.39	0.48	350
19	0.76	0.68	0.72	95
20	0.65	0.56	0.60	71
21	0.74	0.92	0.82	141
22	0.74	0.86	0.80	273
23	0.85	0.88	0.86	249
24	0.65	0.49	0.56	741
25	0.78	0.88	0.83	277

26	0.58	0.72	0.64	67
27	0.79	0.82	0.80	160
28	0.81	0.68	0.74	624
29	0.86	0.89	0.87	249
30	0.76	0.74	0.75	342
31	0.57	0.67	0.61	126
32	0.69	0.83	0.76	125
33	0.57	0.78	0.66	85
34	0.88	0.58	0.70	141
35	0.58	0.40	0.48	62
36	0.80	0.77	0.78	278
37	0.76	0.79	0.77	170
38	0.32	0.20	0.25	90
39	0.85	0.94	0.89	297
40	0.91	0.95	0.93	709
41	0.41	0.19	0.26	57
42	0.29	0.36	0.32	97
43	0.89	0.84	0.86	269
44	0.59	0.36	0.44	76
45	0.70	0.50	0.58	60
46	0.57	0.60	0.58	83
47	0.48	0.28	0.35	421
48	0.42	0.15	0.22	75
49	0.92	0.87	0.89	332
50	0.00	0.00	0.00	79
51	0.49	0.31	0.38	65
52	0.39	0.21	0.27	68
53	0.92	0.91	0.92	411
54	0.40	0.03	0.05	79
55	0.90	0.88	0.89	502
56	0.33	0.34	0.34	89
57	0.48	0.36	0.41	88
58	0.73	0.66	0.69	93
59	0.73	0.42	0.53	107
60	0.41	0.43	0.42	60
61	0.59	0.33	0.42	83
accuracy			0.81	20000
macro avg		0.69	0.66	20000
weighted avg		0.80	0.81	20000

Confusion Matrix:

```
[[795  0   3 ...  0   0   0]
 [  0 919  4 ...  0   0   0]
 [  1   0 959 ...  1   0   2]
 ...
 [  0   2   0 ... 45   1   2]
 [  0   0   0 ...  1  26   0]
 [  0   0  28 ...  1   0  27]]
```

Accuracy: 0.80695

Precision: 0.6939762387610965

Recall: 0.6580846601742629

```
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

## MLP with Hyperparameter tuning

```
In [11]: from sklearn.model_selection import GridSearchCV
from tensorflow.keras.wrappers.scikit_learn import KerasClassifier
# Define the model building function required for KerasClassifier
def create_model(n_hidden_layers=2, n_hidden_neurons=50, activation_function='relu'):
    model = Sequential()
    model.add(Flatten(input_shape=(28*28,)))
    for _ in range(n_hidden_layers):
        model.add(Dense(n_hidden_neurons, activation=activation_function))
    model.add(Dense(len(train_labels[0]), activation='softmax')) # Number of output
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accu
    return model

# Instantiate KerasClassifier with the new model
keras_classifier = KerasClassifier(build_fn=create_model, epochs=10, batch_size=128,
# Define the parameter grid
param_grid = {
    "n_hidden_neurons": [50, 100, 200],
    "activation_function": ["relu", "sigmoid", "tanh"]
}

# Instantiate GridSearchCV
grid = GridSearchCV(estimator=keras_classifier, param_grid=param_grid, cv=3)

# Perform hyperparameter tuning / model selection
grid_result = grid.fit(train_data, train_labels)

# Print the best score and best parameters
print(grid_result.best_score_, grid_result.best_params_)
```

```
C:\Users\Haley\AppData\Local\Temp\ipykernel_10312\3810704986.py:14: DeprecationWarni
ng: KerasClassifier is deprecated, use Sci-Keras (https://github.com/adriangb/scikeras) instead. See https://www.adriangb.com/scikeras/stable/migration.html for help mi
grating.
    keras_classifier = KerasClassifier(build_fn=create_model, epochs=10, batch_size=12
8, verbose=0)
0.8090599973996481 {'activation_function': 'relu', 'n_hidden_neurons': 200}
```

```
In [7]: # Evaluate the model on test data
test_loss, test_acc = grid_result.best_estimator_.model.evaluate(test_data, test_lak
```

```
print('Test accuracy:', test_acc)

# Generate predictions
preds = grid_result.best_estimator_.model.predict(test_data)
preds_classes = np.argmax(preds, axis=1)

# Convert back to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate and print the classification report
classification_report_results = classification_report(test_labels_int, preds_classes)
print("Classification Report:")
print(classification_report_results)

# Generate and print the confusion matrix
confusion_matrix_results = confusion_matrix(test_labels_int, preds_classes)
print("Confusion Matrix:")
print(confusion_matrix_results)

accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro')
recall = recall_score(test_labels_int, preds_classes, average='macro')

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
```



625/625 [=====] - 1s 971us/step - loss: 0.5851 - accuracy: 0.8152

Test accuracy: 0.8151500225067139

625/625 [=====] - 1s 946us/step

Classification Report:

	precision	recall	f1-score	support
0	0.67	0.76	0.72	976
1	0.70	0.82	0.76	1023
2	0.96	0.90	0.92	1003
3	0.94	0.98	0.96	1035
4	0.90	0.96	0.93	903
5	0.86	0.90	0.88	928
6	0.92	0.97	0.94	959
7	0.98	0.97	0.97	1098
8	0.96	0.93	0.94	941
9	0.89	0.98	0.93	929
10	0.85	0.90	0.88	170
11	0.67	0.81	0.73	118
12	0.71	0.88	0.79	316
13	0.87	0.76	0.81	128
14	0.91	0.88	0.90	162
15	0.81	0.79	0.80	261
16	0.65	0.87	0.74	60
17	0.64	0.96	0.77	74
18	0.65	0.36	0.46	350
19	0.76	0.75	0.75	95
20	0.52	0.82	0.63	71
21	0.83	0.93	0.88	141
22	0.76	0.86	0.81	273
23	0.92	0.82	0.87	249
24	0.61	0.59	0.60	741
25	0.76	0.87	0.81	277
26	0.76	0.75	0.75	67
27	0.93	0.78	0.85	160
28	0.79	0.81	0.80	624
29	0.94	0.83	0.88	249
30	0.74	0.92	0.82	342
31	0.62	0.65	0.63	126
32	0.85	0.58	0.69	125
33	0.66	0.60	0.63	85
34	0.76	0.64	0.69	141
35	0.38	0.84	0.52	62
36	0.88	0.75	0.81	278
37	0.93	0.61	0.73	170
38	0.40	0.11	0.17	90
39	0.91	0.92	0.92	297
40	0.92	0.96	0.94	709
41	0.34	0.32	0.33	57
42	0.54	0.21	0.30	97
43	0.91	0.85	0.88	269
44	1.00	0.22	0.37	76
45	0.53	0.67	0.59	60
46	0.51	0.25	0.34	83
47	0.38	0.38	0.38	421
48	0.29	0.12	0.17	75
49	0.91	0.89	0.90	332
50	0.00	0.00	0.00	79
51	0.45	0.26	0.33	65
52	0.44	0.22	0.29	68
53	0.92	0.91	0.92	411
54	0.00	0.00	0.00	79
55	0.79	0.96	0.86	502
56	0.47	0.10	0.17	89

57	0.52	0.34	0.41	88
58	0.59	0.86	0.70	93
59	0.67	0.58	0.62	107
60	0.38	0.25	0.30	60
61	0.54	0.36	0.43	83

accuracy			0.82	20000
macro avg	0.70	0.67	0.67	20000
weighted avg	0.81	0.82	0.80	20000

Confusion Matrix:

```
[[744  0  0 ...  0  0  0]
 [  0 840  2 ...  0  0  0]
 [  1  0 899 ...  0  0 16]
 ...
 [  0  1  0 ... 62  1  4]
 [  0  0  0 ...  0 15  0]
 [  0  0 11 ...  1  0 30]]
```

Accuracy: 0.81515

Precision: 0.6992260869018739

Recall: 0.6688337843850689

## MLP Experiments with varying n\_hidden\_neurons\_values and activation\_function

```
In [12]: #vary n_hidden_neurons_values 10 trials with activation_function='relu'
n_hidden_neurons_values = np.linspace(50, 200, num=10, dtype=int)
n_iterations = 1

best_score = 0
best_params = {}

for n_hidden_neurons in n_hidden_neurons_values:
    for i in range(n_iterations):
        start_time = time.time()

        model = create_model(n_hidden_neurons=n_hidden_neurons, activation_function='relu')
        history = model.fit(train_data, train_labels, epochs=10, batch_size=128, validation_data=(test_data, test_labels))

        # Generate predictions
        preds = model.predict(test_data)
        preds_classes = np.argmax(preds, axis=1)

        # Convert back to integer format
        test_labels_int = np.argmax(test_labels, axis=1)

        accuracy = accuracy_score(test_labels_int, preds_classes)
        precision = precision_score(test_labels_int, preds_classes, average='macro')
        recall = recall_score(test_labels_int, preds_classes, average='macro')

        end_time = time.time()
        elapsed_time = end_time - start_time

        print(f'Iteration {i+1}, n_hidden_neurons: {n_hidden_neurons}')
        print('Accuracy:', accuracy)
        print('Precision:', precision)
        print('Recall:', recall)
        print('Elapsed time:', elapsed_time, 'seconds')
```

```
        if accuracy > best_score:
            best_score = accuracy
            best_params = {'n_hidden_neurons': n_hidden_neurons}

print('Best parameters:', best_params, 'with highest accuracy:', best_score)
```

625/625 [=====] - 0s 673us/step

Iteration 1, n\_hidden\_neurons: 50

Accuracy: 0.78675

Precision: 0.6636166882125001

Recall: 0.6185079213946566

Elapsed time: 12.830229997634888 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 784us/step

Iteration 1, n\_hidden\_neurons: 66

Accuracy: 0.79215

Precision: 0.665171854438555

Recall: 0.634355040173728

Elapsed time: 13.400078296661377 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 795us/step

Iteration 1, n\_hidden\_neurons: 83

Accuracy: 0.80295

Precision: 0.6812350108740798

Recall: 0.6428924862446574

Elapsed time: 13.961905241012573 seconds

625/625 [=====] - 1s 956us/step

Iteration 1, n\_hidden\_neurons: 100

Accuracy: 0.80245

Precision: 0.6863373685561066

Recall: 0.6486132174422466

Elapsed time: 17.535341262817383 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 746us/step

Iteration 1, n\_hidden\_neurons: 116

Accuracy: 0.80855

Precision: 0.693858153776578

Recall: 0.6528442668639657

Elapsed time: 18.015787363052368 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 1ms/step

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

```

Iteration 1, n_hidden_neurons: 133
Accuracy: 0.8125
Precision: 0.6981965852330315
Recall: 0.6641398118730842
Elapsed time: 28.224685192108154 seconds
625/625 [=====] - 1s 1ms/step
Iteration 1, n_hidden_neurons: 150
Accuracy: 0.81245
Precision: 0.701345135716705
Recall: 0.6544264604049093
Elapsed time: 32.15014863014221 seconds
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
625/625 [=====] - 1s 1ms/step
Iteration 1, n_hidden_neurons: 166
Accuracy: 0.8147
Precision: 0.6929582315307035
Recall: 0.6727570945268735
Elapsed time: 25.518718242645264 seconds
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
625/625 [=====] - 1s 1ms/step
Iteration 1, n_hidden_neurons: 183
Accuracy: 0.8154
Precision: 0.6939548768308034
Recall: 0.6648246034446986
Elapsed time: 25.49815011024475 seconds
625/625 [=====] - 1s 1ms/step
Iteration 1, n_hidden_neurons: 200
Accuracy: 0.8126
Precision: 0.7020516458116924
Recall: 0.6721542609509867
Elapsed time: 26.66808271408081 seconds
Best parameters: {'n_hidden_neurons': 183} with highest accuracy: 0.8154
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

```

In [13]: #vary n_hidden_neurons_values 10 trials with activation_function='sigmoid'
n_hidden_neurons_values = np.linspace(50, 200, num=10, dtype=int)
n_iterations = 1

best_score = 0
best_params = {}

for n_hidden_neurons in n_hidden_neurons_values:
    for i in range(n_iterations):
        start_time = time.time()

        model = create_model(n_hidden_neurons=n_hidden_neurons, activation_function=
        history = model.fit(train_data, train_labels, epochs=10, batch_size=128, val

        # Generate predictions
        preds = model.predict(test_data)
        preds_classes = np.argmax(preds, axis=1)

        # Convert back to integer format
        test_labels_int = np.argmax(test_labels, axis=1)

```

```

accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro')
recall = recall_score(test_labels_int, preds_classes, average='macro')

end_time = time.time()
elapsed_time = end_time - start_time

print(f'Iteration {i+1}, n_hidden_neurons: {n_hidden_neurons}')
print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('Elapsed time:', elapsed_time, 'seconds')

if accuracy > best_score:
    best_score = accuracy
    best_params = {'n_hidden_neurons': n_hidden_neurons}

print('Best parameters:', best_params, 'with highest accuracy:', best_score)

```

625/625 [=====] - 1s 850us/step

Iteration 1, n\_hidden\_neurons: 50

Accuracy: 0.75675

Precision: 0.6096183510782207

Recall: 0.5539012216941204

Elapsed time: 16.84760856628418 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
 no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 760us/step

Iteration 1, n\_hidden\_neurons: 66

Accuracy: 0.76995

Precision: 0.6250326141981106

Recall: 0.5751443126918576

Elapsed time: 16.511566877365112 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
 no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 747us/step

Iteration 1, n\_hidden\_neurons: 83

Accuracy: 0.78395

Precision: 0.6599860845206879

Recall: 0.599468787446257

Elapsed time: 16.601797342300415 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
 no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 889us/step

Iteration 1, n\_hidden\_neurons: 100

Accuracy: 0.7901

Precision: 0.6757952152924077

Recall: 0.6182493251699354

Elapsed time: 18.88583779335022 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
 no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 1ms/step

Iteration 1, n\_hidden\_neurons: 116

Accuracy: 0.79835

Precision: 0.6875708204894133

Recall: 0.6297269322283052

Elapsed time: 20.639202117919922 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 818us/step

Iteration 1, n\_hidden\_neurons: 133

Accuracy: 0.80285

Precision: 0.6873893240335153

Recall: 0.6400336714468474

Elapsed time: 19.75173306465149 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 1ms/step

Iteration 1, n\_hidden\_neurons: 150

Accuracy: 0.8053

Precision: 0.68870807018782

Recall: 0.6350585191371019

Elapsed time: 22.533874034881592 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 979us/step

Iteration 1, n\_hidden\_neurons: 166

Accuracy: 0.81055

Precision: 0.702978613200429

Recall: 0.6487298083366403

Elapsed time: 19.90397357940674 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 863us/step

Iteration 1, n\_hidden\_neurons: 183

Accuracy: 0.80925

Precision: 0.6879398439059695

Recall: 0.6429953671238094

Elapsed time: 19.76059055328369 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 1ms/step

Iteration 1, n\_hidden\_neurons: 200

Accuracy: 0.81155

Precision: 0.699349335494755

Recall: 0.6540769711868644

Elapsed time: 31.357375860214233 seconds

Best parameters: {'n\_hidden\_neurons': 200} with highest accuracy: 0.81155

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

```

In [117... #vary n_hidden_neurons_values 10 trials with activation_function='tanh'
n_hidden_neurons_values = np.linspace(50, 200, num=10, dtype=int)
n_iterations = 1

best_score = 0
best_params = {}

for n_hidden_neurons in n_hidden_neurons_values:
    for i in range(n_iterations):
        start_time = time.time()

        model = create_model(n_hidden_neurons=n_hidden_neurons, activation_function=
        history = model.fit(train_data, train_labels, epochs=10, batch_size=128, val

        # Generate predictions
        preds = model.predict(test_data)
        preds_classes = np.argmax(preds, axis=1)

        # Convert back to integer format
        test_labels_int = np.argmax(test_labels, axis=1)

        accuracy = accuracy_score(test_labels_int, preds_classes)
        precision = precision_score(test_labels_int, preds_classes, average='macro')
        recall = recall_score(test_labels_int, preds_classes, average='macro')

        end_time = time.time()
        elapsed_time = end_time - start_time

        print(f'Iteration {i+1}, n_hidden_neurons: {n_hidden_neurons}')
        print('Accuracy:', accuracy)
        print('Precision:', precision)
        print('Recall:', recall)
        print('Elapsed time:', elapsed_time, 'seconds')

        if accuracy > best_score:
            best_score = accuracy
            best_params = {'n_hidden_neurons': n_hidden_neurons}

print('Best parameters:', best_params, 'with highest accuracy:', best_score)

```

625/625 [=====] - 0s 726us/step

Iteration 1, n\_hidden\_neurons: 50

Accuracy: 0.78145

Precision: 0.6459470446428905

Recall: 0.6050419646758668

Elapsed time: 12.479187726974487 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
 no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 763us/step

Iteration 1, n\_hidden\_neurons: 66

Accuracy: 0.79365

Precision: 0.6638872597386016

Recall: 0.625086623732739

Elapsed time: 13.616044521331787 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
 no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 748us/step

Iteration 1, n\_hidden\_neurons: 83

Accuracy: 0.7999

Precision: 0.6921451765570633

Recall: 0.6361593499114829

Elapsed time: 16.139023542404175 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 763us/step

Iteration 1, n\_hidden\_neurons: 100

Accuracy: 0.8033

Precision: 0.683118558830502

Recall: 0.6559164557710065

Elapsed time: 15.396340608596802 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 812us/step

Iteration 1, n\_hidden\_neurons: 116

Accuracy: 0.8076

Precision: 0.6872544888282485

Recall: 0.6566628466836997

Elapsed time: 17.027933835983276 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 747us/step

Iteration 1, n\_hidden\_neurons: 133

Accuracy: 0.815

Precision: 0.7000799085887366

Recall: 0.661848493854145

Elapsed time: 16.12827205657959 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 816us/step

Iteration 1, n\_hidden\_neurons: 150

Accuracy: 0.80685

Precision: 0.6904438187115848

Recall: 0.6547353161267282

Elapsed time: 17.236314058303833 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

625/625 [=====] - 1s 817us/step

Iteration 1, n\_hidden\_neurons: 166

Accuracy: 0.81215

Precision: 0.6900095237354827

Recall: 0.6676738368377416

Elapsed time: 15.4945068359375 seconds

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))



```

625/625 [=====] - 1s 1ms/step
Iteration 1, n_hidden_neurons: 183
Accuracy: 0.8155
Precision: 0.7003523276479232
Recall: 0.6663832058986104
Elapsed time: 24.31060528755188 seconds
625/625 [=====] - 1s 1ms/step
Iteration 1, n_hidden_neurons: 200
Accuracy: 0.80835
Precision: 0.7131020251544583
Recall: 0.664274854221722
Elapsed time: 26.549508571624756 seconds
Best parameters: {'n_hidden_neurons': 183} with highest accuracy: 0.8155
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## CNN before tuning

rerun the first block and continue

```

In [17]: import numpy as np
import pickle
import time
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten
from tensorflow.keras.utils import to_categorical

# Reshape the data for CNN
train_data = train_data.reshape((-1, 28, 28, 1))
test_data = test_data.reshape((-1, 28, 28, 1))

# Normalize the data
train_data = train_data / 255.0
test_data = test_data / 255.0

# Convert the labels to one-hot encoding
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

# Create the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(len(train_labels[0]), activation='softmax') # Number of output classes
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Record the start time
start_time = time.time()

```

```
# Train the model
model.fit(train_data, train_labels, epochs=10, batch_size=128, validation_split=0.1)

# Record the end time
end_time = time.time()

# Calculate and print the training time
train_time = end_time - start_time
print('Training time:', train_time)

# Evaluate the model on test data
test_loss, test_acc = model.evaluate(test_data, test_labels)
print('Test accuracy:', test_acc)

# Generate predictions
preds = model.predict(test_data)
preds_classes = np.argmax(preds, axis=1)

# Convert back to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate and print the classification report
classification_report_results = classification_report(test_labels_int, preds_classes)
print("Classification Report:")
print(classification_report_results)

# Generate and print the confusion matrix
confusion_matrix_results = confusion_matrix(test_labels_int, preds_classes)
print("Confusion Matrix:")
print(confusion_matrix_results)

accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro')
recall = recall_score(test_labels_int, preds_classes, average='macro')

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
```

Epoch 1/10  
704/704 [=====] - 33s 46ms/step - loss: 0.9589 - accuracy: 0.7249 - val\_loss: 0.6323 - val\_accuracy: 0.7913  
Epoch 2/10  
704/704 [=====] - 37s 52ms/step - loss: 0.5433 - accuracy: 0.8171 - val\_loss: 0.5152 - val\_accuracy: 0.8240  
Epoch 3/10  
704/704 [=====] - 32s 45ms/step - loss: 0.4755 - accuracy: 0.8356 - val\_loss: 0.4770 - val\_accuracy: 0.8405  
Epoch 4/10  
704/704 [=====] - 32s 45ms/step - loss: 0.4358 - accuracy: 0.8463 - val\_loss: 0.4614 - val\_accuracy: 0.8388  
Epoch 5/10  
704/704 [=====] - 32s 45ms/step - loss: 0.4068 - accuracy: 0.8536 - val\_loss: 0.4524 - val\_accuracy: 0.8424  
Epoch 6/10  
704/704 [=====] - 31s 45ms/step - loss: 0.3845 - accuracy: 0.8593 - val\_loss: 0.4441 - val\_accuracy: 0.8439  
Epoch 7/10  
704/704 [=====] - 34s 48ms/step - loss: 0.3643 - accuracy: 0.8652 - val\_loss: 0.4361 - val\_accuracy: 0.8462  
Epoch 8/10  
704/704 [=====] - 35s 49ms/step - loss: 0.3446 - accuracy: 0.8701 - val\_loss: 0.4356 - val\_accuracy: 0.8503  
Epoch 9/10  
704/704 [=====] - 33s 47ms/step - loss: 0.3297 - accuracy: 0.8746 - val\_loss: 0.4406 - val\_accuracy: 0.8449  
Epoch 10/10  
704/704 [=====] - 34s 48ms/step - loss: 0.3119 - accuracy: 0.8811 - val\_loss: 0.4359 - val\_accuracy: 0.8513  
Training time: 332.20822501182556  
625/625 [=====] - 2s 4ms/step - loss: 0.4544 - accuracy: 0.8459  
Test accuracy: 0.8458999991416931  
625/625 [=====] - 2s 3ms/step

#### Classification Report:

	precision	recall	f1-score	support
0	0.66	0.81	0.73	976
1	0.68	0.92	0.78	1023
2	0.92	0.97	0.95	1003
3	0.99	0.99	0.99	1035
4	0.96	0.95	0.96	903
5	0.94	0.92	0.93	928
6	0.97	0.97	0.97	959
7	0.99	0.98	0.98	1098
8	0.98	0.96	0.97	941
9	0.93	0.96	0.95	929
10	0.93	0.94	0.93	170
11	0.83	0.90	0.87	118
12	0.75	0.80	0.77	316
13	0.85	0.88	0.87	128
14	0.93	0.92	0.93	162
15	0.81	0.90	0.86	261
16	0.80	0.88	0.84	60
17	0.86	0.97	0.91	74
18	0.64	0.46	0.54	350
19	0.83	0.73	0.78	95
20	0.62	0.77	0.69	71
21	0.86	0.93	0.89	141
22	0.79	0.81	0.80	273
23	0.92	0.97	0.95	249
24	0.64	0.52	0.58	741
25	0.82	0.87	0.85	277

26	0.85	0.85	0.85	67	
27	0.92	0.89	0.91	160	
28	0.79	0.93	0.86	624	
29	0.90	0.91	0.90	249	
30	0.76	0.89	0.82	342	
31	0.64	0.65	0.64	126	
32	0.75	0.77	0.76	125	
33	0.64	0.69	0.67	85	
34	0.79	0.67	0.73	141	
35	0.62	0.45	0.52	62	
36	0.89	0.85	0.87	278	
37	0.87	0.84	0.85	170	
38	0.28	0.19	0.23	90	
39	0.97	0.96	0.97	297	
40	0.95	0.99	0.97	709	
41	0.40	0.07	0.12	57	
42	0.52	0.52	0.52	97	
43	0.92	0.91	0.91	269	
44	0.67	0.38	0.49	76	
45	0.61	0.57	0.59	60	
46	0.72	0.57	0.64	83	
47	0.54	0.25	0.34	421	
48	0.34	0.28	0.31	75	
49	0.89	0.93	0.91	332	
50	0.00	0.00	0.00	79	
51	0.50	0.37	0.42	65	
52	0.47	0.24	0.31	68	
53	0.95	0.94	0.94	411	
54	0.33	0.01	0.02	79	
55	0.93	0.93	0.93	502	
56	0.37	0.21	0.27	89	
57	0.47	0.40	0.43	88	
58	0.70	0.67	0.69	93	
59	0.70	0.60	0.64	107	
60	0.37	0.42	0.39	60	
61	0.53	0.29	0.38	83	
accuracy				0.85	20000
macro avg		0.73	0.71	0.71	20000
weighted avg		0.83	0.85	0.84	20000

Confusion Matrix:

```
[[793  0  0 ...  0  0  0]
 [  0 937  1 ...  0  0  0]
 [  0  0 976 ...  0  0  3]
 ...
 [  0  0  1 ... 64  0  0]
 [  0  0  0 ...  0 25  0]
 [  0  0 35 ...  1  0 24]]
```

Accuracy: 0.8459

Precision: 0.7332046578940103

Recall: 0.7061215043103424

## CNN with Data Augmentation

```
In [18]: import numpy as np
import pickle
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```

# Data Augmentation
datagen = ImageDataGenerator(
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.1, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total
    height_shift_range=0.1, # randomly shift images vertically (fraction of total h
)
datagen.fit(train_data)

# Create the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Dropout(0.25), # Dropout regularization
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Dropout(0.25), # Dropout regularization
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5), # Dropout regularization
    Dense(len(train_labels[0]), activation='softmax') # Number of output classes
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Record the start time
start_time = time.time()

# Train the model
model.fit(datagen.flow(train_data, train_labels, batch_size=128), epochs=10, validate

# Record the end time
end_time = time.time()

# Calculate and print the training time
train_time = end_time - start_time
print('Training time:', train_time)

# Evaluate the model on test data
test_loss, test_acc = model.evaluate(test_data, test_labels)
print('Test accuracy:', test_acc)

# Generate predictions
preds = model.predict(test_data)
preds_classes = np.argmax(preds, axis=1)

# Convert back to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate and print the classification report
classification_report_results = classification_report(test_labels_int, preds_classes)
print("Classification Report:")
print(classification_report_results)

# Generate and print the confusion matrix
confusion_matrix_results = confusion_matrix(test_labels_int, preds_classes)
print("Confusion Matrix:")
print(confusion_matrix_results)

```

```
accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro')
recall = recall_score(test_labels_int, preds_classes, average='macro')

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
```

Epoch 1/10  
782/782 [=====] - 43s 54ms/step - loss: 1.8065 - accuracy: 0.5191 - val\_loss: 0.6790 - val\_accuracy: 0.7784  
Epoch 2/10  
782/782 [=====] - 46s 58ms/step - loss: 1.0971 - accuracy: 0.6682 - val\_loss: 0.5655 - val\_accuracy: 0.8052  
Epoch 3/10  
782/782 [=====] - 45s 58ms/step - loss: 0.9511 - accuracy: 0.7050 - val\_loss: 0.5258 - val\_accuracy: 0.8152  
Epoch 4/10  
782/782 [=====] - 47s 60ms/step - loss: 0.8669 - accuracy: 0.7268 - val\_loss: 0.4994 - val\_accuracy: 0.8230  
Epoch 5/10  
782/782 [=====] - 44s 57ms/step - loss: 0.8122 - accuracy: 0.7416 - val\_loss: 0.4838 - val\_accuracy: 0.8247  
Epoch 6/10  
782/782 [=====] - 44s 56ms/step - loss: 0.7740 - accuracy: 0.7512 - val\_loss: 0.4706 - val\_accuracy: 0.8328  
Epoch 7/10  
782/782 [=====] - 43s 55ms/step - loss: 0.7516 - accuracy: 0.7572 - val\_loss: 0.4584 - val\_accuracy: 0.8339  
Epoch 8/10  
782/782 [=====] - 44s 56ms/step - loss: 0.7316 - accuracy: 0.7643 - val\_loss: 0.4546 - val\_accuracy: 0.8357  
Epoch 9/10  
782/782 [=====] - 44s 56ms/step - loss: 0.7144 - accuracy: 0.7664 - val\_loss: 0.4485 - val\_accuracy: 0.8362  
Epoch 10/10  
782/782 [=====] - 44s 57ms/step - loss: 0.6953 - accuracy: 0.7740 - val\_loss: 0.4401 - val\_accuracy: 0.8400  
Training time: 444.8043382167816  
625/625 [=====] - 2s 4ms/step - loss: 0.4401 - accuracy: 0.8400  
Test accuracy: 0.8400499820709229  
625/625 [=====] - 2s 3ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.61	0.87	0.72	976
1	0.59	0.98	0.74	1023
2	0.90	0.98	0.94	1003
3	0.99	0.99	0.99	1035
4	0.94	0.98	0.96	903
5	0.88	0.94	0.91	928
6	0.95	0.97	0.96	959
7	0.98	0.99	0.99	1098
8	0.95	0.98	0.97	941
9	0.90	0.97	0.93	929
10	0.94	0.94	0.94	170
11	0.90	0.91	0.90	118
12	0.75	0.94	0.84	316
13	0.85	0.81	0.83	128
14	0.94	0.92	0.93	162
15	0.83	0.94	0.88	261
16	0.90	0.75	0.82	60
17	0.83	0.96	0.89	74
18	0.74	0.27	0.40	350
19	0.82	0.84	0.83	95
20	0.65	0.65	0.65	71
21	0.89	0.94	0.91	141
22	0.76	0.97	0.85	273
23	0.91	0.98	0.94	249
24	0.68	0.40	0.51	741
25	0.80	0.95	0.87	277

26	0.95	0.88	0.91	67
27	0.92	0.95	0.94	160
28	0.80	0.84	0.82	624
29	0.90	0.93	0.91	249
30	0.73	0.93	0.82	342
31	0.56	0.90	0.69	126
32	0.57	0.94	0.71	125
33	0.55	0.87	0.67	85
34	0.74	0.82	0.77	141
35	0.57	0.50	0.53	62
36	0.90	0.83	0.87	278
37	0.89	0.85	0.87	170
38	0.00	0.00	0.00	90
39	0.94	0.96	0.95	297
40	0.96	0.98	0.97	709
41	0.00	0.00	0.00	57
42	0.56	0.21	0.30	97
43	0.96	0.92	0.94	269
44	0.83	0.33	0.47	76
45	0.79	0.52	0.63	60
46	0.70	0.64	0.67	83
47	0.00	0.00	0.00	421
48	0.00	0.00	0.00	75
49	0.95	0.94	0.94	332
50	0.00	0.00	0.00	79
51	0.75	0.05	0.09	65
52	0.80	0.06	0.11	68
53	0.96	0.94	0.95	411
54	0.00	0.00	0.00	79
55	0.93	0.94	0.94	502
56	0.00	0.00	0.00	89
57	0.75	0.03	0.07	88
58	0.76	0.14	0.24	93
59	0.82	0.37	0.51	107
60	0.80	0.07	0.12	60
61	0.75	0.22	0.34	83

accuracy			0.84	20000
macro avg	0.73	0.67	0.66	20000
weighted avg	0.81	0.84	0.81	20000

Confusion Matrix:

```
[[ 851    0    0 ...    0    0    0]
 [    0 1006    4 ...    0    0    0]
 [    1    0  987 ...    0    0    0]
 ...
 [    0    0    1 ...   40    0    0]
 [    0    0    0 ...    0    4    0]
 [    0    0   39 ...    0    0   18]]
```

Accuracy: 0.84005

Precision: 0.72577282007204

Recall: 0.667475164631716



```

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Added more Dropout layers and BatchNormalization layers, EarlyStopping callback, and narrowed the search range of learning rate

rerun the first block and continue

```

In [67]: import numpy as np
import time
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout,
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Reshape and normalize the data
train_data = train_data.reshape((-1, 28, 28, 1)) / 255.0
test_data = test_data.reshape((-1, 28, 28, 1)) / 255.0

# Convert the labels to one-hot encoding
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

# Create the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    BatchNormalization(),
    MaxPooling2D((2, 2)),
    Dropout(0.2),

    Conv2D(64, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D((2, 2)),
    Dropout(0.3),

    Flatten(),

    Dense(128, activation='relu'),

```

```

        BatchNormalization(),
        Dropout(0.4),

        Dense(len(train_labels[0]), activation='softmax') # Number of output classes
    ])

# Set learning rate range
optimizer = Adam(learning_rate=0.001)

# Compile the model
model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

# Create callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=3)
model_checkpoint = ModelCheckpoint('model.h5', save_best_only=True)

# Record the start time
start_time = time.time()

# Train the model
model.fit(train_data, train_labels, epochs=10, batch_size=128, validation_split=0.1,
          callbacks=[early_stopping, model_checkpoint])

# Record the end time
end_time = time.time()

# Calculate and print the training time
train_time = end_time - start_time
print('Training time:', train_time)

# Evaluate the model on test data
test_loss, test_acc = model.evaluate(test_data, test_labels)
print('Test accuracy:', test_acc)

# Generate predictions
preds = model.predict(test_data)
preds_classes = np.argmax(preds, axis=1)

# Convert back to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate and print the classification report
classification_report_results = classification_report(test_labels_int, preds_classes)
print("Classification Report:")
print(classification_report_results)

# Generate and print the confusion matrix
confusion_matrix_results = confusion_matrix(test_labels_int, preds_classes)
print("Confusion Matrix:")
print(confusion_matrix_results)

accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro')
recall = recall_score(test_labels_int, preds_classes, average='macro')

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)

```

Epoch 1/10  
704/704 [=====] - 58s 80ms/step - loss: 1.0874 - accuracy: 0.6948 - val\_loss: 0.6971 - val\_accuracy: 0.7704  
Epoch 2/10  
704/704 [=====] - 60s 86ms/step - loss: 0.6273 - accuracy: 0.7952 - val\_loss: 0.4978 - val\_accuracy: 0.8263  
Epoch 3/10  
704/704 [=====] - 62s 88ms/step - loss: 0.5587 - accuracy: 0.8116 - val\_loss: 0.4623 - val\_accuracy: 0.8423  
Epoch 4/10  
704/704 [=====] - 60s 86ms/step - loss: 0.5230 - accuracy: 0.8214 - val\_loss: 0.4397 - val\_accuracy: 0.8476  
Epoch 5/10  
704/704 [=====] - 53s 76ms/step - loss: 0.5014 - accuracy: 0.8271 - val\_loss: 0.4288 - val\_accuracy: 0.8486  
Epoch 6/10  
704/704 [=====] - 57s 81ms/step - loss: 0.4840 - accuracy: 0.8303 - val\_loss: 0.4183 - val\_accuracy: 0.8508  
Epoch 7/10  
704/704 [=====] - 58s 82ms/step - loss: 0.4727 - accuracy: 0.8348 - val\_loss: 0.4154 - val\_accuracy: 0.8510  
Epoch 8/10  
704/704 [=====] - 58s 83ms/step - loss: 0.4623 - accuracy: 0.8358 - val\_loss: 0.4123 - val\_accuracy: 0.8531  
Epoch 9/10  
704/704 [=====] - 60s 85ms/step - loss: 0.4526 - accuracy: 0.8392 - val\_loss: 0.4106 - val\_accuracy: 0.8556  
Epoch 10/10  
704/704 [=====] - 57s 81ms/step - loss: 0.4447 - accuracy: 0.8411 - val\_loss: 0.4087 - val\_accuracy: 0.8551  
Training time: 583.6247837543488  
625/625 [=====] - 4s 6ms/step - loss: 0.4151 - accuracy: 0.8532  
Test accuracy: 0.8532000184059143  
625/625 [=====] - 3s 5ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.64	0.85	0.73	976
1	0.64	0.96	0.77	1023
2	0.95	0.97	0.96	1003
3	0.99	0.99	0.99	1035
4	0.96	0.97	0.96	903
5	0.93	0.91	0.92	928
6	0.97	0.97	0.97	959
7	0.97	0.99	0.98	1098
8	0.97	0.98	0.97	941
9	0.90	0.98	0.94	929
10	0.93	0.94	0.94	170
11	0.92	0.92	0.92	118
12	0.74	0.96	0.84	316
13	0.87	0.86	0.86	128
14	0.93	0.96	0.95	162
15	0.83	0.93	0.87	261
16	0.83	0.90	0.86	60
17	0.87	1.00	0.93	74
18	0.64	0.43	0.51	350
19	0.81	0.78	0.80	95
20	0.62	0.85	0.71	71
21	0.80	0.94	0.87	141
22	0.77	0.96	0.85	273
23	0.91	0.98	0.95	249
24	0.67	0.47	0.55	741
25	0.82	0.95	0.88	277

26	0.97	0.87	0.91	67
27	0.92	0.94	0.93	160
28	0.79	0.92	0.85	624
29	0.89	0.94	0.91	249
30	0.75	0.94	0.84	342
31	0.64	0.83	0.72	126
32	0.78	0.82	0.80	125
33	0.85	0.46	0.60	85
34	0.73	0.84	0.79	141
35	0.63	0.42	0.50	62
36	0.87	0.88	0.88	278
37	0.90	0.86	0.88	170
38	0.00	0.00	0.00	90
39	0.96	0.97	0.96	297
40	0.97	0.98	0.98	709
41	0.50	0.02	0.03	57
42	0.84	0.27	0.41	97
43	0.95	0.91	0.93	269
44	0.85	0.38	0.53	76
45	0.67	0.58	0.63	60
46	0.81	0.55	0.66	83
47	0.51	0.08	0.14	421
48	0.40	0.03	0.05	75
49	0.94	0.92	0.93	332
50	0.00	0.00	0.00	79
51	0.77	0.15	0.26	65
52	0.75	0.22	0.34	68
53	0.93	0.95	0.94	411
54	0.00	0.00	0.00	79
55	0.93	0.93	0.93	502
56	0.60	0.07	0.12	89
57	0.73	0.27	0.40	88
58	0.80	0.71	0.75	93
59	0.65	0.83	0.73	107
60	0.38	0.10	0.16	60
61	0.57	0.52	0.54	83

accuracy			0.85	20000
macro avg	0.76	0.70	0.70	20000
weighted avg	0.84	0.85	0.83	20000

Confusion Matrix:

```
[[829  0  0 ...  0  0  0]
 [  0 979  0 ...  0  0  1]
 [  1  0 971 ...  1  0  8]
 ...
 [  0  1  0 ... 89  0  1]
 [  0  0  0 ...  0  6  0]
 [  0  0 17 ...  1  0 43]]
```

Accuracy: 0.8532

Precision: 0.7603490179470113

Recall: 0.7026956944244633

```

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## CNN with tuning and solve zero\_division

```

In [77]: from kerastuner.tuners import RandomSearch
        from kerastuner.engine.hyperparameters import HyperParameters

# Define the hyperparameter search space
hyperparameters = HyperParameters()
hyperparameters.Int('conv_1_filters', 32, 128, step=32)
hyperparameters.Float('dropout_1', 0.2, 0.5, step=0.1)
hyperparameters.Int('conv_2_filters', 64, 256, step=32)
hyperparameters.Float('dropout_2', 0.2, 0.5, step=0.1)
hyperparameters.Int('dense_units', 128, 512, step=64)
hyperparameters.Float('dropout_3', 0.2, 0.5, step=0.1)
hyperparameters.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])

# Define the function to build the model
def build_model(hp):
    model = Sequential()
    model.add(Conv2D(hp.Int('conv_1_filters', 32, 128, step=32), (3, 3), activation=
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(hp.Float('dropout_1', 0.2, 0.5, step=0.1)))

    model.add(Conv2D(hp.Int('conv_2_filters', 64, 256, step=32), (3, 3), activation=
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(hp.Float('dropout_2', 0.2, 0.5, step=0.1)))

    model.add(Flatten())

    model.add(Dense(hp.Int('dense_units', 128, 512, step=64), activation='relu'))
    model.add(BatchNormalization())
    model.add(Dropout(hp.Float('dropout_3', 0.2, 0.5, step=0.1)))

    model.add(Dense(len(train_labels[0]), activation='softmax'))

    optimizer = Adam(learning_rate=hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e
    model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['ac

    return model

```

```

# Create the tuner with reduced max_trials
tuner = RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=3,
    executions_per_trial=1,
    directory='tuner_dir',
    project_name='my_model'
)

# Perform hyperparameter tuning with reduced epochs
tuner.search(train_data, train_labels, epochs=3, validation_split=0.1)

# Get the best hyperparameters
best_hps = tuner.get_best_hyperparameters()[0]

# Print the best hyperparameters
print("Best Hyperparameters:")
print(best_hps)

```

Trial 2 Complete [00h 08m 54s]  
val\_accuracy: 0.8285999894142151

Best val\_accuracy So Far: 0.8468999862670898  
Total elapsed time: 00h 08m 54s  
INFO:tensorflow:Oracle triggered exit  
Best Hyperparameters:  
<keras\_tuner.engine.hyperparameters.hyperparameters.HyperParameters object at 0x00000206C8542D00>

```

In [94]: import numpy as np
import time
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout,
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Import Keras Tuner related functions
from kerastuner.tuners import RandomSearch
from kerastuner.engine.hyperparameters import HyperParameters

# Define the hyperparameter search space
hyperparameters = HyperParameters()
hyperparameters.Int('conv_1_filters', 32, 128, step=32)
hyperparameters.Float('dropout_1', 0.2, 0.5, step=0.1)
hyperparameters.Int('conv_2_filters', 64, 256, step=32)
hyperparameters.Float('dropout_2', 0.2, 0.5, step=0.1)
hyperparameters.Int('dense_units', 128, 512, step=64)
hyperparameters.Float('dropout_3', 0.2, 0.5, step=0.1)
hyperparameters.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])

# Assuming train_data, test_data, train_labels, and test_labels are defined

# Reshape and normalize the data
train_data = train_data.reshape((-1, 28, 28, 1)) / 255.0
test_data = test_data.reshape((-1, 28, 28, 1)) / 255.0

# Convert the labels to one-hot encoding
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

# Define the function to build the model

```

```

def build_model(hp):
    model = Sequential()
    model.add(Conv2D(hp.Int('conv_1_filters', 32, 128, step=32), (3, 3), activation=
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(hp.Float('dropout_1', 0.2, 0.5, step=0.1)))

    model.add(Conv2D(hp.Int('conv_2_filters', 64, 256, step=32), (3, 3), activation=
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(hp.Float('dropout_2', 0.2, 0.5, step=0.1)))

    model.add(Flatten())

    model.add(Dense(hp.Int('dense_units', 128, 512, step=64), activation='relu'))
    model.add(BatchNormalization())
    model.add(Dropout(hp.Float('dropout_3', 0.2, 0.5, step=0.1)))

    model.add(Dense(len(train_labels[0]), activation='softmax'))

    optimizer = Adam(learning_rate=hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e
    model.compile(optimizer=optimizer, loss='categorical_crossentropy',
    metrics=['accuracy'])

    return model

# Create the tuner
tuner = RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=3,
    executions_per_trial=1,
    directory='tuner_dir',
    project_name='my_model'
)

# Create callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=3)
model_checkpoint = ModelCheckpoint('model.h5', save_best_only=True)

# Perform hyperparameter tuning
tuner.search(train_data, train_labels, epochs=10, validation_split=0.1,
             callbacks=[early_stopping, model_checkpoint])

# Get the best hyperparameters
best_hps = tuner.get_best_hyperparameters()[0]

# Print the best hyperparameters
print("Best Hyperparameters:")
print(best_hps)

# Build the model with the best hyperparameters
model = build_model(best_hps)

# Train the model
model.fit(train_data, train_labels, epochs=10, batch_size=128, validation_split=0.1,
         callbacks=[early_stopping, model_checkpoint])

# Evaluate the model on test data
test_loss, test_acc = model.evaluate(test_data, test_labels)
print('Test accuracy:', test_acc)

# Generate predictions
preds = model.predict(test_data)

```

```
preds_classes = np.argmax(preds, axis=1)

# Convert back to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate and print the classification report
classification_report_results = classification_report(test_labels_int, preds_classes)
print("Classification Report:")
print(classification_report_results)

# Generate and print the confusion matrix
confusion_matrix_results = confusion_matrix(test_labels_int, preds_classes)
print("Confusion Matrix:")
print(confusion_matrix_results)

# Calculate and print the accuracy, precision, and recall scores
accuracy = accuracy_score(test_labels_int, preds_classes)
precision = precision_score(test_labels_int, preds_classes, average='macro', zero_division=0)
recall = recall_score(test_labels_int, preds_classes, average='macro', zero_division=0)

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
```



```

INFO:tensorflow:Reloading Tuner from tuner_dir\my_model\tuner0.json
INFO:tensorflow:Oracle triggered exit
Best Hyperparameters:
<keras_tuner.engine.hyperparameters.hyperparameters.HyperParameters object at 0x0000
0206D8FB7DF0>
Epoch 1/10
704/704 [=====] - 136s 192ms/step - loss: 1.5989 - accurac
y: 0.6012 - val_loss: 1.2708 - val_accuracy: 0.6839
Epoch 2/10
704/704 [=====] - 137s 195ms/step - loss: 0.9053 - accurac
y: 0.7439 - val_loss: 0.7315 - val_accuracy: 0.7908
Epoch 3/10
704/704 [=====] - 149s 212ms/step - loss: 0.7368 - accurac
y: 0.7795 - val_loss: 0.6266 - val_accuracy: 0.8081
Epoch 4/10
704/704 [=====] - 175s 248ms/step - loss: 0.6520 - accurac
y: 0.7966 - val_loss: 0.5678 - val_accuracy: 0.8220
Epoch 5/10
704/704 [=====] - 190s 269ms/step - loss: 0.5976 - accurac
y: 0.8089 - val_loss: 0.5351 - val_accuracy: 0.8247
Epoch 6/10
704/704 [=====] - 189s 268ms/step - loss: 0.5585 - accurac
y: 0.8166 - val_loss: 0.5039 - val_accuracy: 0.8338
Epoch 7/10
704/704 [=====] - 191s 272ms/step - loss: 0.5290 - accurac
y: 0.8247 - val_loss: 0.4858 - val_accuracy: 0.8355
Epoch 8/10
704/704 [=====] - 202s 287ms/step - loss: 0.5039 - accurac
y: 0.8304 - val_loss: 0.4605 - val_accuracy: 0.8426
Epoch 9/10
704/704 [=====] - 189s 269ms/step - loss: 0.4848 - accurac
y: 0.8359 - val_loss: 0.4487 - val_accuracy: 0.8483
Epoch 10/10
704/704 [=====] - 160s 228ms/step - loss: 0.4697 - accurac
y: 0.8377 - val_loss: 0.4366 - val_accuracy: 0.8515
625/625 [=====] - 7s 12ms/step - loss: 0.4460 - accuracy:
0.8469
Test accuracy: 0.8469499945640564
625/625 [=====] - 7s 11ms/step

```

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.73	0.70	976
1	0.66	0.93	0.77	1023
2	0.93	0.97	0.95	1003
3	0.99	0.99	0.99	1035
4	0.96	0.96	0.96	903
5	0.93	0.91	0.92	928
6	0.94	0.98	0.96	959
7	0.98	0.98	0.98	1098
8	0.93	0.99	0.96	941
9	0.91	0.97	0.94	929
10	0.90	0.95	0.92	170
11	0.88	0.85	0.87	118
12	0.75	0.91	0.82	316
13	0.85	0.84	0.85	128
14	0.94	0.89	0.91	162
15	0.83	0.91	0.86	261
16	0.96	0.80	0.87	60
17	0.91	0.92	0.91	74
18	0.57	0.49	0.53	350
19	0.93	0.67	0.78	95
20	0.63	0.80	0.70	71
21	0.84	0.94	0.89	141

22	0.77	0.95	0.85	273
23	0.89	0.96	0.92	249
24	0.60	0.63	0.62	741
25	0.85	0.86	0.85	277
26	0.90	0.85	0.88	67
27	0.90	0.92	0.91	160
28	0.79	0.92	0.85	624
29	0.91	0.89	0.90	249
30	0.77	0.90	0.83	342
31	0.61	0.75	0.67	126
32	0.80	0.78	0.79	125
33	0.67	0.71	0.69	85
34	0.73	0.79	0.76	141
35	0.65	0.21	0.32	62
36	0.86	0.86	0.86	278
37	0.91	0.81	0.86	170
38	0.00	0.00	0.00	90
39	0.94	0.97	0.95	297
40	0.94	0.99	0.96	709
41	0.43	0.05	0.09	57
42	0.69	0.26	0.38	97
43	0.94	0.90	0.92	269
44	0.77	0.36	0.49	76
45	0.73	0.68	0.71	60
46	0.75	0.58	0.65	83
47	0.56	0.15	0.23	421
48	0.43	0.04	0.07	75
49	0.92	0.93	0.92	332
50	1.00	0.00	0.00	79
51	0.47	0.46	0.47	65
52	0.74	0.21	0.32	68
53	0.90	0.95	0.93	411
54	1.00	0.00	0.00	79
55	0.93	0.93	0.93	502
56	0.50	0.20	0.29	89
57	0.64	0.31	0.42	88
58	0.78	0.74	0.76	93
59	0.77	0.64	0.70	107
60	0.43	0.20	0.27	60
61	0.65	0.48	0.55	83

accuracy			0.85	20000
macro avg	0.78	0.70	0.70	20000
weighted avg	0.84	0.85	0.83	20000

Confusion Matrix:

```
[[712  0  1 ...  0  0  0]
 [  0 949  2 ...  0  0  0]
 [  0  0 973 ...  1  0  2]
 ...
 [  0  1  1 ... 69  0  0]
 [  0  0  0 ...  0 12  0]
 [  0  0 25 ...  1  0 40]]
```

Accuracy: 0.84695

Precision: 0.7808151860624953

Recall: 0.6966060166276296

## SVM Experiments with varying C and Kernels

```

In [6]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.model_selection import cross_val_score, train_test_split
import numpy as np
import time

# Define parameter grid
C_values = np.logspace(-3, 2, num=5)

# Define kernel types
kernels = ['rbf', 'poly']

# Split your data so that you only use 50% of it for training
sample_X_train, _, sample_y_train, _ = train_test_split(X_train, y_train, test_size=0.5)

best_score = 0
best_params = {}

# Perform grid search on the training data for each kernel type and 'C' value
for kernel in kernels:
    for C in C_values:
        start = time.time()

        # Create an instance of the SVM model with current kernel
        svm_model_param_grid = SVC(kernel=kernel, C=C)

        # Perform cross-validation on the training data
        scores = cross_val_score(svm_model_param_grid, sample_X_train, sample_y_train, cv=5)

        # Compute the mean accuracy and check if it's the best score so far
        mean_score = np.mean(scores)
        if mean_score > best_score:
            best_score = mean_score
            best_params = {'kernel': kernel, 'C': C}

        # Fit the model with the 'C' parameter
        svm_model_param_grid.fit(sample_X_train, sample_y_train)

        # Make predictions on the test set using the model
        predictions_SVMBestparam = svm_model_param_grid.predict(X_test)

        # Calculate evaluation metrics
        accuracy = accuracy_score(y_test, predictions_SVMBestparam)
        precision = precision_score(y_test, predictions_SVMBestparam, average='macro')
        recall = recall_score(y_test, predictions_SVMBestparam, average='macro')
        confusion_mat = confusion_matrix(y_test, predictions_SVMBestparam)

        end = time.time()

        print("Parameters: Kernel=", kernel, "C=", C)
        print("Mean cross-validation accuracy:", mean_score)
        print("Accuracy:", accuracy)
        print("Precision:", precision)
        print("Recall:", recall)
        print("Elapsed Time:", end - start)
        print()

print("Best Parameters:", best_params, "with highest mean cross-validation accuracy:")

```

```

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

Parameters: Kernel= rbf C= 0.001  
Mean cross-validation accuracy: 0.05535000006918577  
Accuracy: 0.05465  
Precision: 0.007791782423387355  
Recall: 0.016175291432800838  
Elapsed Time: 5008.530836582184

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

Parameters: Kernel= rbf C= 0.01778279410038923  
Mean cross-validation accuracy: 0.4880500087348441  
Accuracy: 0.533  
Precision: 0.22705525025965226  
Recall: 0.20633334240237858  
Elapsed Time: 22268.483466863632

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

Parameters: Kernel= rbf C= 0.31622776601683794  
Mean cross-validation accuracy: 0.7150249908940088  
Accuracy: 0.73145  
Precision: 0.5956326855218941  
Recall: 0.47159356019462234  
Elapsed Time: 1928.0591492652893

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

Parameters: Kernel= rbf C= 5.623413251903491  
Mean cross-validation accuracy: 0.7782250147224131  
Accuracy: 0.78765  
Precision: 0.6843231429789959  
Recall: 0.5996393949271969  
Elapsed Time: 1553.9620718955994

Parameters: Kernel= rbf C= 100.0  
Mean cross-validation accuracy: 0.7664500203325542  
Accuracy: 0.77735  
Precision: 0.6745260369351598  
Recall: 0.601002881319926  
Elapsed Time: 1818.9180636405945

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

Parameters: Kernel= poly C= 0.001  
Mean cross-validation accuracy: 0.05562499944454513  
Accuracy: 0.0551  
Precision: 0.047407157440332165  
Recall: 0.017502732863971263  
Elapsed Time: 8902.655396938324

C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318:  
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with  
no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

```
Parameters: Kernel= poly C= 0.01778279410038923
Mean cross-validation accuracy: 0.1416752476709023
Accuracy: 0.26
Precision: 0.36394742064784563
Recall: 0.10906894234203014
Elapsed Time: 3895.9232897758484
```

```
C:\Users\Haley\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with
no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
Parameters: Kernel= poly C= 0.31622776601683794
Mean cross-validation accuracy: 0.6480001283067924
Accuracy: 0.6809
Precision: 0.5754524787241456
Recall: 0.3982451637096467
Elapsed Time: 2705.840261220932
```

```
Parameters: Kernel= poly C= 5.623413251903491
Mean cross-validation accuracy: 0.7730000715894604
Accuracy: 0.7869
Precision: 0.6780583818063879
Recall: 0.5922134056978122
Elapsed Time: 1857.6277301311493
```

```
Parameters: Kernel= poly C= 100.0
Mean cross-validation accuracy: 0.7689250353352729
Accuracy: 0.78105
Precision: 0.6726165518699603
Recall: 0.6062399787837118
Elapsed Time: 1790.865995168686
```

```
Best Parameters: {'kernel': 'rbf', 'C': 5.623413251903491} with highest mean cross-v
alidation accuracy: 0.7782250147224131
```

In [ ]:

# Linear model

```
In [1]: #Import all necessary libraries
import numpy as np
import pickle
import matplotlib.pyplot as plt
import time
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.model_selection import GridSearchCV

# Load the emnist_train.pkl file
with open('/kaggle/input/emnist-trainpkl/emnist_train.pkl', 'rb') as f:
    train_dict = pickle.load(f, encoding='bytes')

# Extract the data and labels arrays from the dictionary
train_data = train_dict['data']
train_labels = train_dict['labels']

# Load the emnist_test.pkl file
with open('/kaggle/input/emnist-testpkl/emnist_test.pkl', 'rb') as f:
    test_dict = pickle.load(f, encoding='bytes')

# Extract the data and labels arrays from the dictionary
test_data = test_dict['data']
test_labels = test_dict['labels']

# Convert the data and labels arrays to NumPy arrays
train_data = np.array(train_data, dtype=np.float32)
train_labels = np.array(train_labels, dtype=np.int32)
test_data = np.array(test_data, dtype=np.float32)
test_labels = np.array(test_labels, dtype=np.int32)

# Define class names
# Create class names for all 62 classes
class_names = ['Class {}'.format(i) for i in range(62)]

#Plot a grid of EMNIST examples of a specified size."""

def plot_examples(data, pred, target, n_rows=3, n_cols=7):
    """Plot a grid of EMNIST examples of a specified size."""
    # Size figure depending on the size of the grid
    plt.figure(figsize=(n_cols * 1.2, n_rows * 2.0))

    for row in range(n_rows):
        for col in range(n_cols):
            # Get the next index of the image
            index = n_cols * row + col

            # Reshape the flattened image to its original shape
            image = data[index].reshape((28, 28))

            # Plot the image at the appropriate place in the grid
```

```

        plt.subplot(n_rows, n_cols, index + 1)
        plt.imshow(image, cmap="binary")
        plt.axis('off')
        plt.title(class_names[target[index]] + '\n' + '>>' + class_names[pred[index]])

plt.tight_layout()
plt.show()


# Plot examples from different classes
fig, axes = plt.subplots(8, 8, figsize=(12, 12))
for i, ax in enumerate(axes.flat):
    if i < len(class_names):
        class_index = i # Compute the class index
        # Get samples of the current class
        class_samples = train_data[train_labels == class_index]
        # Randomly select a sample from the class
        sample_index = np.random.randint(class_samples.shape[0])
        ax.imshow(class_samples[sample_index].reshape((28, 28)), cmap='gray')
        ax.axis('off')
        # Set the title as the class name
        ax.set_title(class_names[class_index])
    else:
        ax.axis('off')

plt.tight_layout()
plt.show()

```

/opt/conda/lib/python3.10/site-packages/scipy/\_\_init\_\_.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.23.5

warnings.warn(f"A NumPy version >={np\_minversion} and <{np\_maxversion}")



```
In [2]: #Data preprocessing
#Perform simple Normalise and reshape data
train_images = train_data / 255.0
test_images = test_data / 255.0

# Reshape the data from (28, 28) to (784,)
train_images = train_data.reshape(-1, 784)
test_images = test_data.reshape(-1, 784)

# Normalize the data by standardization
scaler = StandardScaler()
train_images = scaler.fit_transform(train_images)
test_images = scaler.transform(test_images)
```

```
In [3]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(train_images, train_labels, test
```

## SVM with Linear kernel

```
In [ ]: # Create an instance of the linear SVM model
svm_model = SVC(kernel='linear')

start = time.time()
```



```

# Train the model
svm_model.fit(X_train, y_train)

end = time.time()
print(end - start)

# Make predictions on the test set
predictions = svm_model.predict(X_test)

start = time.time()

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, predictions)
precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
confusion_mat = confusion_matrix(y_test, predictions)

end = time.time()
print(end - start)

# Print the evaluation metrics and confusion matrix
print("Accuracy_Linear:", accuracy)
print("Precision_Linear:", precision)
print("Recall_Linear:", recall)
print("Confusion Matrix_Linear:")
print(confusion_mat)

# Plot examples of linear SVM
plot_examples(X_test, predictions, y_test, n_rows=3, n_cols=7)

```

978.5045220851898

0.02817225456237793

Accuracy\_Linear: 0.7264

Precision\_Linear: 0.581932100433202


Recall\_Linear: 0.5626831399061308

Confusion Matrix\_Linear:

```

[[ 731    0    0 ...    0    0    0]
 [   0 1005    2 ...    0    0    0]
 [   2    3  837 ...    0    0   20]
 ...
 [   0    3    0 ...   39    2    0]
 [   0    3    1 ...    1   21    0]
 [   0    0   35 ...    0    1   24]]

```



## PCA compression on Linear model

```
In [ ]: # apply PCA without reducing dimensionality, then compute the min number of dimension
pca = PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
d = np.argmax(cumsum >= 0.95) + 1

# Apply PCA for dimensionality reduction

pca = PCA(n_components=d)
X_train_pca = pca.fit_transform(X_train)
# Transform the test set
X_test_pca = pca.transform(X_test)

# Fit the SVM model on the transformed training set
svm_model = SVC(kernel='linear')

start = time.time()

svm_model.fit(X_train_pca, y_train)

end = time.time()
print(end - start)

# Make predictions on the transformed test set
predictions_pca = svm_model.predict(X_test_pca)

start = time.time()
# Calculate accuracy
accuracy = accuracy_score(y_test, predictions_pca)

end = time.time()
print(end - start)

print("Reduced data Accuracy:", accuracy)
```

```
# Plotting the reduced data to find if they form uniform clusters

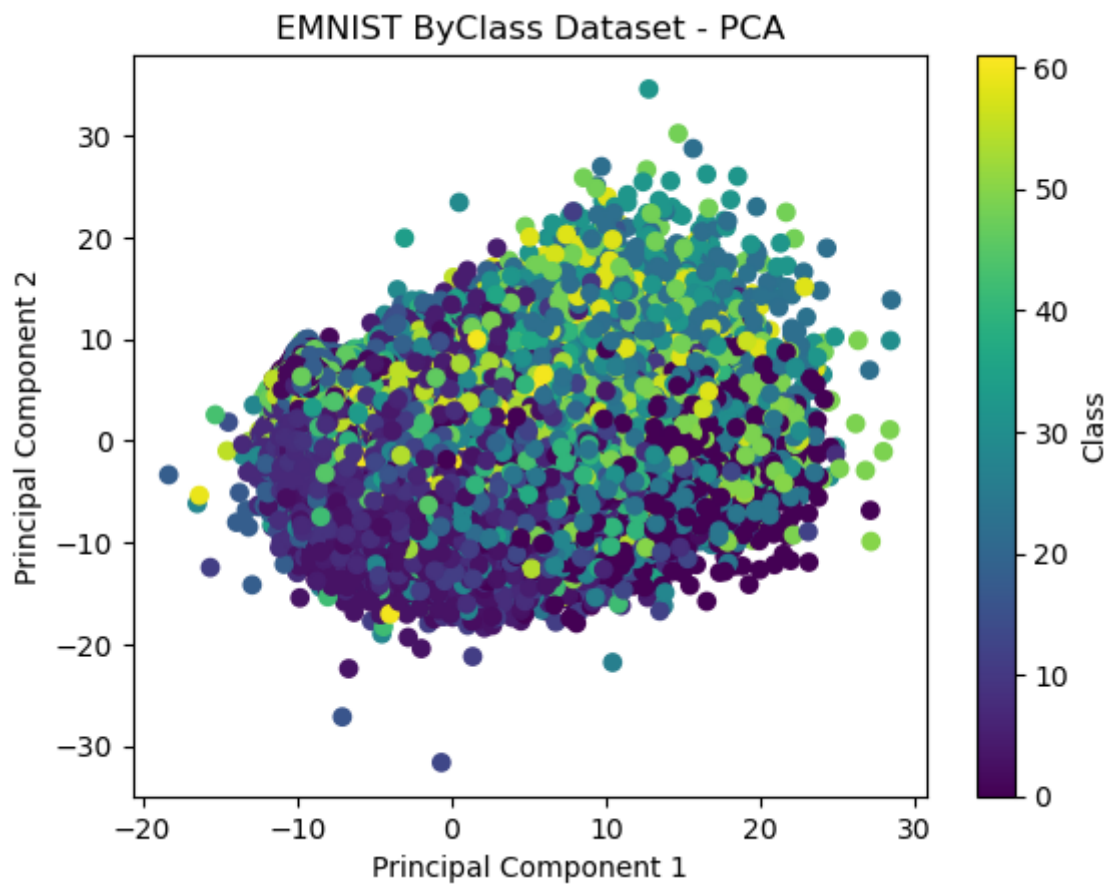
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train)
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('EMNIST ByClass Dataset - PCA')
plt.colorbar(label='Class')
plt.savefig('pca.pdf')
plt.show()


# Plot examples of PCA compressed SVM results
plot_examples(X_test, predictions_pca, y_test, n_rows=5, n_cols=5)
```

843.0075702667236

0.0025217533111572266

Reduced data Accuracy: 0.7494





SVM with default paramters

```
In [ ]: # Create an instance of the SVM model with default hyperparameters
svm_model = SVC(kernel='poly', C=1.0, gamma='scale')

start = time.time()
# Train the model
svm_model.fit(X_train, y_train)
end = time.time()
print(end - start)

# Make predictions on the test set
predictions_SVMD = svm_model.predict(X_test)

start = time.time()

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, predictions_SVMD)
precision = precision_score(y_test, predictions_SVMD, average='macro')
recall = recall_score(y_test, predictions_SVMD, average='macro')
confusion_mat = confusion_matrix(y_test, predictions_SVMD)

end = time.time()
print(end - start)

# Print the evaluation metrics and confusion matrix
print("Accuracy_SVMD:", accuracy)
print("Precision_SVMD:", precision)
print("Recall_SVMD:", recall)
print("Confusion Matrix_SVMD:")
print(confusion_mat)
```

1925.7620568275452

0.03128767013549805

Accuracy\_SVMD: 0.78095

Precision\_SVMD: 0.6770057099086038

Recall\_SVMD: 0.5627796728995796

Confusion Matrix\_SVMD:

```
[[ 832    0    1 ...    0    0    0]
 [   0 1051    1 ...    0    0    0]
 [   1    0  928 ...    0    0    7]
 ...
 [   0    0    2 ...   45    0    0]
 [   0    3    1 ...    0    1    0]
 [   0    0  50 ...    0    0  21]]
```

/suphys/mlan9395/miniconda3/lib/python3.10/site-packages/sklearn/metrics/\_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

# SVM with hyperparameter tuning

```
In [7]: # Define the parameter grid
param_grid = {
    'C': [ 0.1, 1.0, 10, 100],
    'gamma': ['scale', 'auto']
}

# Create an instance of the SVM model
svm_model_param_grid = SVC(kernel='rbf')

# Split the data into training and testing sets for gridsearch
X_train_Grid, X_test_Grid, y_train_Grid, y_test_Grid = train_test_split(train_images

# Create an instance of the GridSearchCV with the SVM model and parameter grid
grid_search = GridSearchCV(svm_model_param_grid, param_grid, cv=5, scoring='accuracy

start = time.time()
# Perform grid search on the training data
grid_search.fit(X_train_Grid, y_train_Grid)
end = time.time()
print(end - start)

# Get the best parameters and best model
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

# Make predictions on the test set using the best model
predictions_SVMBestparam = best_model.predict(X_test_Grid)

start = time.time()
# Calculate evaluation metrics
accuracy = accuracy_score(y_test_Grid, predictions_SVMBestparam)
precision = precision_score(y_test_Grid, predictions_SVMBestparam, average='macro')
recall = recall_score(y_test_Grid, predictions_SVMBestparam, average='macro')
confusion_mat = confusion_matrix(y_test_Grid, predictions_SVMBestparam)

end = time.time()
print(end - start)

# Print the evaluation metrics and confusion matrix
print("Best Parameters:", best_params)
print("Accuracy_SVMBestparam:", accuracy)
print("Precision_SVMBestparam:", precision)
print("Recall_SVMBestparam:", recall)
print("Confusion Matrix_SVMBestparam:")
print(confusion_mat)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

```

/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A NumPy
version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version
1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A NumPy
version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version
1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A NumPy
version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version
1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A NumPy
version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version
1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
2376.185211658478
0.14995598793029785
Best Parameters: {'C': 10, 'gamma': 'auto'}
Accuracy_SVMBestparam: 0.766675
Precision_SVMBestparam: 0.651683821326552
Recall_SVMBestparam: 0.5713692727767514
Confusion Matrix_SVMBestparam:
[[2920    1    1 ...    0    0    0]
 [   0 4166    9 ...    0    0    0]
 [   12    4 3624 ...    2    0   32]
 ...
 [   0    1    5 ...  190    1    0]
 [   0    5    2 ...    1   38    0]
 [   0    0  146 ...    0    0   92]]

```

### SVM with best parameters

```

In [9]: # Create an instance of the SVM model with the best hyperparameters
svm_model = SVC(kernel='rbf', C=10, gamma='auto')

start = time.time()
svm_model.fit(X_train, y_train)
end = time.time()
print(end - start)

# Make predictions
predictions_SVMNew = svm_model.predict(X_test)

start = time.time()
# Calculate evaluation metrics
accuracy = accuracy_score(y_test, predictions_SVMNew)
precision = precision_score(y_test, predictions_SVMNew, average='macro')
recall = recall_score(y_test, predictions_SVMNew, average='macro')
confusion_mat = confusion_matrix(y_test, predictions_SVMNew)
end = time.time()
print(end - start)

# Print the evaluation metrics and confusion matrix
print("Accuracy_SVMNew:", accuracy)
print("Precision_SVMNew:", precision)
print("Recall_SVMNew:", recall)
print("Confusion Matrix_SVMNew:")
print(confusion_mat)


# Plot examples of best parameterized SVM
plot_examples(X_test, predictions_SVMNew, y_test, n_rows=3, n_cols=7)

```

```

1378.204790353775
0.0479891300201416
Accuracy_SVMNew: 0.8042
Precision_SVMNew: 0.7041671297346184
Recall_SVMNew: 0.6339506453348212
Confusion Matrix_SVMNew:
[[ 777    0    1 ...    0    0    0]
 [    0 1029    0 ...    0    0    0]
 [    2    2  913 ...    0    0    8]
 ...
 [    0    0    1 ...   53    1    0]
 [    0    0    1 ...    0   16    0]
 [    0    0   29 ...    0    0   39]]

```



## SVM with Regularization

```

In [10]: # Access the value of the C parameter
best_C = best_params['C']

# Apply regularization by reducing C value
regularized_C = best_C * 0.1 # Example: reduce C by 10 times

# Update the C parameter in the best_params dictionary
best_params['C'] = regularized_C

# Initialize a new SVM classifier with the regularized C value
regularized_svm = SVC(**best_params)

start = time.time()

# Fit the regularized SVM on the training data
regularized_svm.fit(X_train, y_train)

end = time.time()
print(end - start)

# Make predictions on the test set using the regularized model
predictions-Regularized = regularized_svm.predict(X_test)

```



```

start = time.time()
# Calculate evaluation metrics
accuracy = accuracy_score(y_test, predictions-Regularized)
precision = precision_score(y_test, predictions-Regularized, average='macro')
recall = recall_score(y_test, predictions-Regularized, average='macro')
confusion_mat = confusion_matrix(y_test, predictions-Regularized)

end = time.time()
print(end - start)

# Print the evaluation metrics and confusion matrix
print("Regularized Parameters:", best_params)
print("Accuracy-Regularized:", accuracy)
print("Precision-Regularized:", precision)
print("Recall-Regularized:", recall)
print("Confusion Matrix-Regularized:")
print(confusion_mat)

# Plot examples of Regularized SVM
plot_examples(X_test, predictions-Regularized, y_test, n_rows=3, n_cols=7)

```

1454.2729988098145

/opt/conda/lib/python3.10/site-packages/sklearn/metrics/\_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
warn\_prf(average, modifier, msg\_start, len(result))

0.04742550849914551

Regularized Parameters: {'C': 1.0, 'gamma': 'auto'}

Accuracy-Regularized: 0.79505

Precision-Regularized: 0.6775295893378837


Recall-Regularized: 0.6022630503219472

Confusion Matrix-Regularized:

```

[[ 830   0   0 ...   0   0   0]
 [   0 1036   0 ...   0   0   0]
 [   3   2  919 ...   0   0   7]
 ...
 [   0   0   1 ...  54   1   0]
 [   0   3   1 ...   1   8   0]
 [   0   0  42 ...   0   0  26]]

```





# CNN Experiments with varying learning\_rates and dense\_units\_options

```
In [3]: import numpy as np
import time
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout,
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Reshape and normalize the data
train_data = train_data.reshape((-1, 28, 28, 1)) / 255.0
test_data = test_data.reshape((-1, 28, 28, 1)) / 255.0

# Convert the labels to one-hot encoding
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

# Define the function to build the model
def build_model(units, learning_rate):
    model = Sequential()
    model.add(Conv2D(64, (3, 3), activation='relu', input_shape=(28, 28, 1)))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(0.3))

    model.add(Conv2D(128, (3, 3), activation='relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(0.3))

    model.add(Flatten())

    model.add(Dense(units, activation='relu'))
    model.add(BatchNormalization())
    model.add(Dropout(0.3))

    model.add(Dense(len(train_labels[0]), activation='softmax'))

    optimizer = Adam(learning_rate=learning_rate)
    model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['ac

    return model

# Define the parameter grid
dense_units_options = list(range(128, 512 + 1, 128)) # 10 values from 128 to 512
learning_rates = [1e-2, 1e-3, 1e-4, 1e-5] # 4 values for learning rate

# Create callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=3)
model_checkpoint = ModelCheckpoint('model.h5', save_best_only=True)

# Record results
results = []

# Perform manual hyperparameter tuning
for units in dense_units_options:
```

```

for learning_rate in learning_rates:
    start_time = time.time()

    model = build_model(units, learning_rate)

    model.fit(train_data, train_labels, epochs=10, batch_size=128, validation_split=0.1,
              callbacks=[early_stopping, model_checkpoint])

    test_loss, test_acc = model.evaluate(test_data, test_labels)
    end_time = time.time()

    # Generate predictions
    preds = model.predict(test_data)
    preds_classes = np.argmax(preds, axis=1)

    # Convert back to integer format
    test_labels_int = np.argmax(test_labels, axis=1)

    accuracy = accuracy_score(test_labels_int, preds_classes)
    precision = precision_score(test_labels_int, preds_classes, average='macro', zero_division=0)
    recall = recall_score(test_labels_int, preds_classes, average='macro', zero_division=0)

    results.append({
        'learning_rate': learning_rate,
        'units': units,
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'runtime': end_time - start_time
    })

    print("Learning rate:", learning_rate)
    print("Units:", units)
    print("Accuracy:", accuracy)
    print("Precision:", precision)
    print("Recall:", recall)
    print("Runtime:", end_time - start_time)
    print()

# Print final results
for result in results:
    print("Learning rate:", result['learning_rate'])
    print("Units:", result['units'])
    print("Accuracy:", result['accuracy'])
    print("Precision:", result['precision'])
    print("Recall:", result['recall'])
    print("Runtime:", result['runtime'])
    print()

```

Epoch 1/10  
704/704 [=====] - 94s 131ms/step - loss: 0.7751 - accuracy:  
0.7564 - val\_loss: 0.5234 - val\_accuracy: 0.8188  
Epoch 2/10  
704/704 [=====] - 90s 128ms/step - loss: 0.5641 - accuracy:  
0.8077 - val\_loss: 0.4837 - val\_accuracy: 0.8300  
Epoch 3/10  
704/704 [=====] - 90s 128ms/step - loss: 0.5172 - accuracy:  
0.8206 - val\_loss: 0.4549 - val\_accuracy: 0.8425  
Epoch 4/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4919 - accuracy:  
0.8268 - val\_loss: 0.4275 - val\_accuracy: 0.8488  
Epoch 5/10  
704/704 [=====] - 92s 131ms/step - loss: 0.4751 - accuracy:  
0.8324 - val\_loss: 0.4386 - val\_accuracy: 0.8447  
Epoch 6/10  
704/704 [=====] - 91s 130ms/step - loss: 0.4601 - accuracy:  
0.8361 - val\_loss: 0.4170 - val\_accuracy: 0.8531  
Epoch 7/10  
704/704 [=====] - 92s 130ms/step - loss: 0.4534 - accuracy:  
0.8368 - val\_loss: 0.4127 - val\_accuracy: 0.8506  
Epoch 8/10  
704/704 [=====] - 92s 131ms/step - loss: 0.4396 - accuracy:  
0.8410 - val\_loss: 0.4221 - val\_accuracy: 0.8507  
Epoch 9/10  
704/704 [=====] - 90s 128ms/step - loss: 0.4317 - accuracy:  
0.8431 - val\_loss: 0.4083 - val\_accuracy: 0.8535  
Epoch 10/10  
704/704 [=====] - 89s 127ms/step - loss: 0.4262 - accuracy:  
0.8460 - val\_loss: 0.4129 - val\_accuracy: 0.8529  
625/625 [=====] - 7s 11ms/step - loss: 0.4304 - accuracy:  
0.8477  
625/625 [=====] - 7s 11ms/step  
Learning rate: 0.01  
Units: 128  
Accuracy: 0.8477  
Precision: 0.8043123489373756  
Recall: 0.6988400193387927  
Runtime: 919.8214240074158

Epoch 1/10  
704/704 [=====] - 92s 129ms/step - loss: 0.9372 - accuracy:  
0.7319 - val\_loss: 0.6508 - val\_accuracy: 0.7973  
Epoch 2/10  
704/704 [=====] - 91s 129ms/step - loss: 0.5620 - accuracy:  
0.8119 - val\_loss: 0.4706 - val\_accuracy: 0.8392  
Epoch 3/10  
704/704 [=====] - 90s 128ms/step - loss: 0.5042 - accuracy:  
0.8271 - val\_loss: 0.4387 - val\_accuracy: 0.8479  
Epoch 4/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4696 - accuracy:  
0.8368 - val\_loss: 0.4292 - val\_accuracy: 0.8491  
Epoch 5/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4490 - accuracy:  
0.8414 - val\_loss: 0.4260 - val\_accuracy: 0.8506  
Epoch 6/10  
704/704 [=====] - 91s 130ms/step - loss: 0.4352 - accuracy:  
0.8447 - val\_loss: 0.4094 - val\_accuracy: 0.8575  
Epoch 7/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4191 - accuracy:  
0.8489 - val\_loss: 0.4011 - val\_accuracy: 0.8548  
Epoch 8/10  
704/704 [=====] - 92s 130ms/step - loss: 0.4114 - accuracy:  
0.8498 - val\_loss: 0.4013 - val\_accuracy: 0.8601

Epoch 9/10  
704/704 [=====] - 90s 128ms/step - loss: 0.4024 - accuracy:  
0.8539 - val\_loss: 0.3975 - val\_accuracy: 0.8568  
Epoch 10/10  
704/704 [=====] - 89s 127ms/step - loss: 0.3928 - accuracy:  
0.8559 - val\_loss: 0.4042 - val\_accuracy: 0.8557  
625/625 [=====] - 7s 11ms/step - loss: 0.4190 - accuracy:  
0.8526  
625/625 [=====] - 7s 11ms/step  
Learning rate: 0.001  
Units: 128  
Accuracy: 0.85255  
Precision: 0.7867328400171182  
Recall: 0.707327582008939  
Runtime: 915.5359442234039

Epoch 1/10  
704/704 [=====] - 90s 126ms/step - loss: 1.8558 - accuracy:  
0.5451 - val\_loss: 1.3337 - val\_accuracy: 0.6647  
Epoch 2/10  
704/704 [=====] - 90s 127ms/step - loss: 1.0449 - accuracy:  
0.7092 - val\_loss: 0.7781 - val\_accuracy: 0.7750  
Epoch 3/10  
704/704 [=====] - 89s 126ms/step - loss: 0.8387 - accuracy:  
0.7549 - val\_loss: 0.6669 - val\_accuracy: 0.8006  
Epoch 4/10  
704/704 [=====] - 90s 127ms/step - loss: 0.7355 - accuracy:  
0.7750 - val\_loss: 0.5861 - val\_accuracy: 0.8153  
Epoch 5/10  
704/704 [=====] - 90s 127ms/step - loss: 0.6693 - accuracy:  
0.7913 - val\_loss: 0.5480 - val\_accuracy: 0.8231  
Epoch 6/10  
704/704 [=====] - 91s 129ms/step - loss: 0.6227 - accuracy:  
0.8007 - val\_loss: 0.5232 - val\_accuracy: 0.8296  
Epoch 7/10  
704/704 [=====] - 90s 128ms/step - loss: 0.5872 - accuracy:  
0.8103 - val\_loss: 0.5001 - val\_accuracy: 0.8354  
Epoch 8/10  
704/704 [=====] - 90s 127ms/step - loss: 0.5607 - accuracy:  
0.8161 - val\_loss: 0.4794 - val\_accuracy: 0.8405  
Epoch 9/10  
704/704 [=====] - 88s 126ms/step - loss: 0.5394 - accuracy:  
0.8214 - val\_loss: 0.4703 - val\_accuracy: 0.8414  
Epoch 10/10  
704/704 [=====] - 89s 126ms/step - loss: 0.5197 - accuracy:  
0.8265 - val\_loss: 0.4558 - val\_accuracy: 0.8434  
625/625 [=====] - 7s 11ms/step - loss: 0.4596 - accuracy:  
0.8430  
625/625 [=====] - 7s 11ms/step  
Learning rate: 0.0001  
Units: 128  
Accuracy: 0.843  
Precision: 0.786125959157098  
Recall: 0.6869673416696267  
Runtime: 901.746609210968

Epoch 1/10  
704/704 [=====] - 89s 125ms/step - loss: 3.5610 - accuracy:  
0.2011 - val\_loss: 2.5119 - val\_accuracy: 0.4362  
Epoch 2/10  
704/704 [=====] - 86s 123ms/step - loss: 2.3527 - accuracy:  
0.4388 - val\_loss: 1.8504 - val\_accuracy: 0.5708  
Epoch 3/10  
704/704 [=====] - 87s 124ms/step - loss: 1.9694 - accuracy:

0.5211 - val\_loss: 1.6330 - val\_accuracy: 0.6141  
Epoch 4/10  
704/704 [=====] - 89s 126ms/step - loss: 1.7571 - accuracy:  
0.5642 - val\_loss: 1.4801 - val\_accuracy: 0.6429  
Epoch 5/10  
704/704 [=====] - 87s 124ms/step - loss: 1.5993 - accuracy:  
0.5968 - val\_loss: 1.3559 - val\_accuracy: 0.6660  
Epoch 6/10  
704/704 [=====] - 88s 125ms/step - loss: 1.4837 - accuracy:  
0.6179 - val\_loss: 1.2523 - val\_accuracy: 0.6847  
Epoch 7/10  
704/704 [=====] - 88s 125ms/step - loss: 1.3779 - accuracy:  
0.6399 - val\_loss: 1.1683 - val\_accuracy: 0.7006  
Epoch 8/10  
704/704 [=====] - 89s 126ms/step - loss: 1.2928 - accuracy:  
0.6607 - val\_loss: 1.0958 - val\_accuracy: 0.7164  
Epoch 9/10  
704/704 [=====] - 88s 126ms/step - loss: 1.2194 - accuracy:  
0.6743 - val\_loss: 1.0361 - val\_accuracy: 0.7265  
Epoch 10/10  
704/704 [=====] - 87s 124ms/step - loss: 1.1573 - accuracy:  
0.6891 - val\_loss: 0.9822 - val\_accuracy: 0.7369  
625/625 [=====] - 6s 10ms/step - loss: 0.9743 - accuracy:  
0.7347  
625/625 [=====] - 7s 10ms/step  
Learning rate: 1e-05  
Units: 128  
Accuracy: 0.73465  
Precision: 0.7162887549989162  
Recall: 0.47558125347175484  
Runtime: 885.5641613006592

Epoch 1/10  
704/704 [=====] - 94s 132ms/step - loss: 0.7820 - accuracy:  
0.7578 - val\_loss: 0.5630 - val\_accuracy: 0.8104  
Epoch 2/10  
704/704 [=====] - 91s 130ms/step - loss: 0.5575 - accuracy:  
0.8092 - val\_loss: 0.4729 - val\_accuracy: 0.8359  
Epoch 3/10  
704/704 [=====] - 92s 130ms/step - loss: 0.5061 - accuracy:  
0.8226 - val\_loss: 0.4778 - val\_accuracy: 0.8396  
Epoch 4/10  
704/704 [=====] - 96s 136ms/step - loss: 0.4787 - accuracy:  
0.8310 - val\_loss: 0.4482 - val\_accuracy: 0.8374  
Epoch 5/10  
704/704 [=====] - 94s 133ms/step - loss: 0.4574 - accuracy:  
0.8370 - val\_loss: 0.4365 - val\_accuracy: 0.8489  
Epoch 6/10  
704/704 [=====] - 92s 130ms/step - loss: 0.4473 - accuracy:  
0.8396 - val\_loss: 0.4465 - val\_accuracy: 0.8459  
Epoch 7/10  
704/704 [=====] - 93s 132ms/step - loss: 0.4364 - accuracy:  
0.8432 - val\_loss: 0.4151 - val\_accuracy: 0.8531  
Epoch 8/10  
704/704 [=====] - 92s 131ms/step - loss: 0.4266 - accuracy:  
0.8447 - val\_loss: 0.4369 - val\_accuracy: 0.8388  
Epoch 9/10  
704/704 [=====] - 91s 130ms/step - loss: 0.4209 - accuracy:  
0.8456 - val\_loss: 0.4305 - val\_accuracy: 0.8474  
Epoch 10/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4103 - accuracy:  
0.8486 - val\_loss: 0.4248 - val\_accuracy: 0.8481  
625/625 [=====] - 7s 11ms/step - loss: 0.4298 - accuracy:  
0.8505

625/625 [=====] - 7s 11ms/step  
Learning rate: 0.01  
Units: 256  
Accuracy: 0.85055  
Precision: 0.8134135377605697  
Recall: 0.7079299811099503  
Runtime: 990.635219335556

Epoch 1/10  
704/704 [=====] - 92s 129ms/step - loss: 0.8462 - accuracy:  
0.7467 - val\_loss: 0.5616 - val\_accuracy: 0.8180  
Epoch 2/10  
704/704 [=====] - 89s 127ms/step - loss: 0.5305 - accuracy:  
0.8183 - val\_loss: 0.4596 - val\_accuracy: 0.8417  
Epoch 3/10  
704/704 [=====] - 90s 128ms/step - loss: 0.4730 - accuracy:  
0.8350 - val\_loss: 0.4297 - val\_accuracy: 0.8502  
Epoch 4/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4441 - accuracy:  
0.8429 - val\_loss: 0.4255 - val\_accuracy: 0.8490  
Epoch 5/10  
704/704 [=====] - 91s 129ms/step - loss: 0.4249 - accuracy:  
0.8463 - val\_loss: 0.4127 - val\_accuracy: 0.8448  
Epoch 6/10  
704/704 [=====] - 91s 130ms/step - loss: 0.4065 - accuracy:  
0.8515 - val\_loss: 0.3999 - val\_accuracy: 0.8518  
Epoch 7/10  
704/704 [=====] - 91s 129ms/step - loss: 0.3911 - accuracy:  
0.8558 - val\_loss: 0.3967 - val\_accuracy: 0.8589  
Epoch 8/10  
704/704 [=====] - 90s 128ms/step - loss: 0.3816 - accuracy:  
0.8582 - val\_loss: 0.4018 - val\_accuracy: 0.8547  
Epoch 9/10  
704/704 [=====] - 89s 126ms/step - loss: 0.3693 - accuracy:  
0.8619 - val\_loss: 0.3934 - val\_accuracy: 0.8603  
Epoch 10/10  
704/704 [=====] - 89s 126ms/step - loss: 0.3607 - accuracy:  
0.8642 - val\_loss: 0.3876 - val\_accuracy: 0.8601  
625/625 [=====] - 7s 11ms/step - loss: 0.4048 - accuracy:  
0.8572  
625/625 [=====] - 7s 11ms/step  
Learning rate: 0.001  
Units: 256  
Accuracy: 0.85725  
Precision: 0.7893143274662201  
Recall: 0.7183029748170656  
Runtime: 909.7056584358215

Epoch 1/10  
704/704 [=====] - 92s 129ms/step - loss: 1.6217 - accuracy:  
0.5851 - val\_loss: 1.0392 - val\_accuracy: 0.7184  
Epoch 2/10  
704/704 [=====] - 91s 129ms/step - loss: 0.8919 - accuracy:  
0.7355 - val\_loss: 0.6605 - val\_accuracy: 0.7935  
Epoch 3/10  
704/704 [=====] - 90s 128ms/step - loss: 0.7346 - accuracy:  
0.7727 - val\_loss: 0.5750 - val\_accuracy: 0.8150  
Epoch 4/10  
704/704 [=====] - 90s 128ms/step - loss: 0.6565 - accuracy:  
0.7895 - val\_loss: 0.5346 - val\_accuracy: 0.8231  
Epoch 5/10  
704/704 [=====] - 91s 129ms/step - loss: 0.6066 - accuracy:  
0.8029 - val\_loss: 0.5024 - val\_accuracy: 0.8307  
Epoch 6/10



704/704 [=====] - 90s 129ms/step - loss: 0.5661 - accuracy:  
0.8131 - val\_loss: 0.4779 - val\_accuracy: 0.8378  
Epoch 7/10  
704/704 [=====] - 90s 128ms/step - loss: 0.5386 - accuracy:  
0.8196 - val\_loss: 0.4647 - val\_accuracy: 0.8417  
Epoch 8/10  
704/704 [=====] - 90s 127ms/step - loss: 0.5176 - accuracy:  
0.8254 - val\_loss: 0.4527 - val\_accuracy: 0.8441  
Epoch 9/10  
704/704 [=====] - 90s 127ms/step - loss: 0.4966 - accuracy:  
0.8299 - val\_loss: 0.4425 - val\_accuracy: 0.8454  
Epoch 10/10  
704/704 [=====] - 90s 128ms/step - loss: 0.4800 - accuracy:  
0.8344 - val\_loss: 0.4327 - val\_accuracy: 0.8488  
625/625 [=====] - 7s 11ms/step - loss: 0.4431 - accuracy:  
0.8464  
625/625 [=====] - 7s 11ms/step  
Learning rate: 0.0001  
Units: 256  
Accuracy: 0.8464  
Precision: 0.7865026115391711  
Recall: 0.6944006943451094  
Runtime: 911.2518379688263

Epoch 1/10  
704/704 [=====] - 93s 129ms/step - loss: 3.4790 - accuracy:  
0.2348 - val\_loss: 2.2663 - val\_accuracy: 0.4781  
Epoch 2/10  
704/704 [=====] - 92s 130ms/step - loss: 2.1198 - accuracy:  
0.4841 - val\_loss: 1.5505 - val\_accuracy: 0.6142  
Epoch 3/10  
704/704 [=====] - 92s 131ms/step - loss: 1.7303 - accuracy:  
0.5613 - val\_loss: 1.3260 - val\_accuracy: 0.6611  
Epoch 4/10  
704/704 [=====] - 92s 130ms/step - loss: 1.5093 - accuracy:  
0.6047 - val\_loss: 1.1733 - val\_accuracy: 0.6919  
Epoch 5/10  
704/704 [=====] - 93s 132ms/step - loss: 1.3563 - accuracy:  
0.6343 - val\_loss: 1.0655 - val\_accuracy: 0.7136  
Epoch 6/10  
704/704 [=====] - 102s 144ms/step - loss: 1.2381 - accuracy:  
0.6615 - val\_loss: 0.9814 - val\_accuracy: 0.7302  
Epoch 7/10  
704/704 [=====] - 93s 131ms/step - loss: 1.1483 - accuracy:  
0.6790 - val\_loss: 0.9133 - val\_accuracy: 0.7439  
Epoch 8/10  
704/704 [=====] - 94s 134ms/step - loss: 1.0787 - accuracy:  
0.6953 - val\_loss: 0.8616 - val\_accuracy: 0.7541  
Epoch 9/10  
704/704 [=====] - 96s 136ms/step - loss: 1.0179 - accuracy:  
0.7081 - val\_loss: 0.8148 - val\_accuracy: 0.7647  
Epoch 10/10  
704/704 [=====] - 96s 136ms/step - loss: 0.9673 - accuracy:  
0.7183 - val\_loss: 0.7792 - val\_accuracy: 0.7728  
625/625 [=====] - 8s 12ms/step - loss: 0.7850 - accuracy:  
0.7670  
625/625 [=====] - 7s 12ms/step  
Learning rate: 1e-05  
Units: 256  
Accuracy: 0.76695  
Precision: 0.6844942561186504  
Recall: 0.5549426168871804  
Runtime: 991.9178223609924

Epoch 1/10  
704/704 [=====] - 98s 137ms/step - loss: 0.7771 - accuracy:  
0.7574 - val\_loss: 0.5857 - val\_accuracy: 0.7930  
Epoch 2/10  
704/704 [=====] - 96s 136ms/step - loss: 0.5504 - accuracy:  
0.8120 - val\_loss: 0.7374 - val\_accuracy: 0.8183  
Epoch 3/10  
704/704 [=====] - 95s 135ms/step - loss: 0.4989 - accuracy:  
0.8262 - val\_loss: 0.4714 - val\_accuracy: 0.8364  
Epoch 4/10  
704/704 [=====] - 96s 136ms/step - loss: 0.4784 - accuracy:  
0.8307 - val\_loss: 0.4543 - val\_accuracy: 0.8425  
Epoch 5/10  
704/704 [=====] - 95s 136ms/step - loss: 0.4600 - accuracy:  
0.8362 - val\_loss: 0.4374 - val\_accuracy: 0.8482  
Epoch 6/10  
704/704 [=====] - 95s 134ms/step - loss: 0.4468 - accuracy:  
0.8390 - val\_loss: 0.4356 - val\_accuracy: 0.8438  
Epoch 7/10  
704/704 [=====] - 94s 134ms/step - loss: 0.4355 - accuracy:  
0.8420 - val\_loss: 0.4349 - val\_accuracy: 0.8498  
Epoch 8/10  
704/704 [=====] - 94s 134ms/step - loss: 0.4244 - accuracy:  
0.8453 - val\_loss: 0.4450 - val\_accuracy: 0.8421  
Epoch 9/10  
704/704 [=====] - 95s 135ms/step - loss: 0.4175 - accuracy:  
0.8481 - val\_loss: 0.4446 - val\_accuracy: 0.8392  
Epoch 10/10  
704/704 [=====] - 95s 135ms/step - loss: 0.4109 - accuracy:  
0.8473 - val\_loss: 0.4226 - val\_accuracy: 0.8521  
625/625 [=====] - 8s 13ms/step - loss: 0.4446 - accuracy:  
0.8467  
625/625 [=====] - 8s 12ms/step  
Learning rate: 0.01  
Units: 384  
Accuracy: 0.84665  
Precision: 0.7905247884057616  
Recall: 0.7130944499335667  
Runtime: 991.4726612567902

Epoch 1/10  
704/704 [=====] - 97s 136ms/step - loss: 0.8197 - accuracy:  
0.7496 - val\_loss: 0.5699 - val\_accuracy: 0.8108  
Epoch 2/10  
704/704 [=====] - 95s 135ms/step - loss: 0.5246 - accuracy:  
0.8209 - val\_loss: 0.4502 - val\_accuracy: 0.8384  
Epoch 3/10  
704/704 [=====] - 96s 136ms/step - loss: 0.4687 - accuracy:  
0.8347 - val\_loss: 0.4225 - val\_accuracy: 0.8515  
Epoch 4/10  
704/704 [=====] - 96s 136ms/step - loss: 0.4366 - accuracy:  
0.8433 - val\_loss: 0.4177 - val\_accuracy: 0.8516  
Epoch 5/10  
704/704 [=====] - 97s 138ms/step - loss: 0.4172 - accuracy:  
0.8486 - val\_loss: 0.4084 - val\_accuracy: 0.8535  
Epoch 6/10  
704/704 [=====] - 96s 137ms/step - loss: 0.4008 - accuracy:  
0.8532 - val\_loss: 0.4102 - val\_accuracy: 0.8524  
Epoch 7/10  
704/704 [=====] - 97s 137ms/step - loss: 0.3893 - accuracy:  
0.8546 - val\_loss: 0.4039 - val\_accuracy: 0.8553  
Epoch 8/10  
704/704 [=====] - 96s 136ms/step - loss: 0.3742 - accuracy:  
0.8595 - val\_loss: 0.3915 - val\_accuracy: 0.8603

Epoch 9/10  
704/704 [=====] - 97s 138ms/step - loss: 0.3593 - accuracy: 0.8649 - val\_loss: 0.4111 - val\_accuracy: 0.8559  
Epoch 10/10  
704/704 [=====] - 106s 151ms/step - loss: 0.3541 - accuracy: 0.8663 - val\_loss: 0.4045 - val\_accuracy: 0.8574  
625/625 [=====] - 9s 15ms/step - loss: 0.4177 - accuracy: 0.8541  
625/625 [=====] - 9s 14ms/step  
Learning rate: 0.001  
Units: 384  
Accuracy: 0.8541  
Precision: 0.7530003563189905  
Recall: 0.7165125146157162  
Runtime: 984.3839774131775

Epoch 1/10  
704/704 [=====] - 114s 160ms/step - loss: 1.5209 - accuracy: 0.6029 - val\_loss: 0.9681 - val\_accuracy: 0.7277  
Epoch 2/10  
704/704 [=====] - 100s 142ms/step - loss: 0.8376 - accuracy: 0.7451 - val\_loss: 0.6143 - val\_accuracy: 0.8062  
Epoch 3/10  
704/704 [=====] - 99s 141ms/step - loss: 0.6872 - accuracy: 0.7821 - val\_loss: 0.5401 - val\_accuracy: 0.8230  
Epoch 4/10  
704/704 [=====] - 100s 142ms/step - loss: 0.6162 - accuracy: 0.7989 - val\_loss: 0.5034 - val\_accuracy: 0.8319  
Epoch 5/10  
704/704 [=====] - 99s 140ms/step - loss: 0.5718 - accuracy: 0.8111 - val\_loss: 0.4765 - val\_accuracy: 0.8386  
Epoch 6/10  
704/704 [=====] - 99s 140ms/step - loss: 0.5344 - accuracy: 0.8194 - val\_loss: 0.4610 - val\_accuracy: 0.8420  
Epoch 7/10  
704/704 [=====] - 99s 140ms/step - loss: 0.5117 - accuracy: 0.8252 - val\_loss: 0.4497 - val\_accuracy: 0.8454  
Epoch 8/10  
704/704 [=====] - 96s 137ms/step - loss: 0.4883 - accuracy: 0.8315 - val\_loss: 0.4391 - val\_accuracy: 0.8487  
Epoch 9/10  
704/704 [=====] - 97s 137ms/step - loss: 0.4730 - accuracy: 0.8366 - val\_loss: 0.4291 - val\_accuracy: 0.8523  
Epoch 10/10  
704/704 [=====] - 104s 148ms/step - loss: 0.4580 - accuracy: 0.8405 - val\_loss: 0.4310 - val\_accuracy: 0.8452  
625/625 [=====] - 9s 14ms/step - loss: 0.4393 - accuracy: 0.8433  
625/625 [=====] - 8s 13ms/step  
Learning rate: 0.0001  
Units: 384  
Accuracy: 0.8433  
Precision: 0.7891048209621495  
Recall: 0.7014753587907294  
Runtime: 1054.4654290676117

Epoch 1/10  
704/704 [=====] - 113s 158ms/step - loss: 3.1625 - accuracy: 0.2915 - val\_loss: 2.0095 - val\_accuracy: 0.5307  
Epoch 2/10  
704/704 [=====] - 105s 149ms/step - loss: 1.9249 - accuracy: 0.5191 - val\_loss: 1.3929 - val\_accuracy: 0.6357  
Epoch 3/10  
704/704 [=====] - 104s 147ms/step - loss: 1.5803 - accuracy:

y: 0.5829 - val\_loss: 1.1975 - val\_accuracy: 0.6724  
Epoch 4/10  
704/704 [=====] - 105s 149ms/step - loss: 1.3831 - accuracy: 0.6237 - val\_loss: 1.0631 - val\_accuracy: 0.7012  
Epoch 5/10  
704/704 [=====] - 104s 148ms/step - loss: 1.2509 - accuracy: 0.6508 - val\_loss: 0.9703 - val\_accuracy: 0.7195  
Epoch 6/10  
704/704 [=====] - 107s 153ms/step - loss: 1.1451 - accuracy: 0.6761 - val\_loss: 0.8917 - val\_accuracy: 0.7382  
Epoch 7/10  
704/704 [=====] - 105s 149ms/step - loss: 1.0663 - accuracy: 0.6927 - val\_loss: 0.8434 - val\_accuracy: 0.7503  
Epoch 8/10  
704/704 [=====] - 106s 150ms/step - loss: 1.0006 - accuracy: 0.7074 - val\_loss: 0.7892 - val\_accuracy: 0.7617  
Epoch 9/10  
704/704 [=====] - 104s 148ms/step - loss: 0.9516 - accuracy: 0.7181 - val\_loss: 0.7535 - val\_accuracy: 0.7695  
Epoch 10/10  
704/704 [=====] - 103s 147ms/step - loss: 0.9097 - accuracy: 0.7285 - val\_loss: 0.7191 - val\_accuracy: 0.7776  
625/625 [=====] - 8s 14ms/step - loss: 0.7206 - accuracy: 0.7772  
625/625 [=====] - 8s 13ms/step  
Learning rate: 1e-05  
Units: 384  
Accuracy: 0.77725  
Precision: 0.690018898958171  
Recall: 0.5786048612733493  
Runtime: 1114.2524206638336

Epoch 1/10  
704/704 [=====] - 111s 156ms/step - loss: 0.8052 - accuracy: 0.7523 - val\_loss: 0.5326 - val\_accuracy: 0.8195  
Epoch 2/10  
704/704 [=====] - 108s 153ms/step - loss: 0.5536 - accuracy: 0.8114 - val\_loss: 0.5141 - val\_accuracy: 0.8290  
Epoch 3/10  
704/704 [=====] - 109s 155ms/step - loss: 0.5092 - accuracy: 0.8212 - val\_loss: 0.5487 - val\_accuracy: 0.8251  
Epoch 4/10  
704/704 [=====] - 109s 154ms/step - loss: 0.4825 - accuracy: 0.8295 - val\_loss: 0.4652 - val\_accuracy: 0.8374  
Epoch 5/10  
704/704 [=====] - 109s 155ms/step - loss: 0.4670 - accuracy: 0.8330 - val\_loss: 0.4528 - val\_accuracy: 0.8490  
Epoch 6/10  
704/704 [=====] - 107s 153ms/step - loss: 0.4557 - accuracy: 0.8369 - val\_loss: 0.4435 - val\_accuracy: 0.8449  
Epoch 7/10  
704/704 [=====] - 108s 153ms/step - loss: 0.4455 - accuracy: 0.8388 - val\_loss: 0.4561 - val\_accuracy: 0.8357  
Epoch 8/10  
704/704 [=====] - 110s 156ms/step - loss: 0.4346 - accuracy: 0.8416 - val\_loss: 0.4329 - val\_accuracy: 0.8476  
Epoch 9/10  
704/704 [=====] - 108s 153ms/step - loss: 0.4245 - accuracy: 0.8450 - val\_loss: 0.4492 - val\_accuracy: 0.8448  
Epoch 10/10  
704/704 [=====] - 108s 154ms/step - loss: 0.4203 - accuracy: 0.8458 - val\_loss: 0.4303 - val\_accuracy: 0.8485  
625/625 [=====] - 9s 15ms/step - loss: 0.4465 - accuracy: 0.8453

625/625 [=====] - 9s 14ms/step  
Learning rate: 0.01  
Units: 512  
Accuracy: 0.84525  
Precision: 0.7942784310891914  
Recall: 0.7128416009952518  
Runtime: 1098.0348672866821

Epoch 1/10  
704/704 [=====] - 113s 158ms/step - loss: 0.8159 - accuracy: 0.7522 - val\_loss: 0.6532 - val\_accuracy: 0.7742  
Epoch 2/10  
704/704 [=====] - 109s 155ms/step - loss: 0.5232 - accuracy: 0.8205 - val\_loss: 0.4610 - val\_accuracy: 0.8402  
Epoch 3/10  
704/704 [=====] - 109s 154ms/step - loss: 0.4649 - accuracy: 0.8348 - val\_loss: 0.4390 - val\_accuracy: 0.8397  
Epoch 4/10  
704/704 [=====] - 108s 154ms/step - loss: 0.4367 - accuracy: 0.8422 - val\_loss: 0.4354 - val\_accuracy: 0.8441  
Epoch 5/10  
704/704 [=====] - 111s 158ms/step - loss: 0.4122 - accuracy: 0.8492 - val\_loss: 0.4147 - val\_accuracy: 0.8561  
Epoch 6/10  
704/704 [=====] - 107s 152ms/step - loss: 0.3977 - accuracy: 0.8527 - val\_loss: 0.4080 - val\_accuracy: 0.8558  
Epoch 7/10  
704/704 [=====] - 107s 152ms/step - loss: 0.3819 - accuracy: 0.8573 - val\_loss: 0.4029 - val\_accuracy: 0.8512  
Epoch 8/10  
704/704 [=====] - 109s 154ms/step - loss: 0.3658 - accuracy: 0.8619 - val\_loss: 0.3970 - val\_accuracy: 0.8589  
Epoch 9/10  
704/704 [=====] - 107s 151ms/step - loss: 0.3554 - accuracy: 0.8648 - val\_loss: 0.3986 - val\_accuracy: 0.8603  
Epoch 10/10  
704/704 [=====] - 105s 150ms/step - loss: 0.3428 - accuracy: 0.8687 - val\_loss: 0.4071 - val\_accuracy: 0.8555  
625/625 [=====] - 9s 14ms/step - loss: 0.4220 - accuracy: 0.8522  
625/625 [=====] - 9s 14ms/step  
Learning rate: 0.001  
Units: 512  
Accuracy: 0.85215  
Precision: 0.7637378673224909  
Recall: 0.7207188433737077  
Runtime: 1113.2304723262787

Epoch 1/10  
704/704 [=====] - 108s 152ms/step - loss: 1.4339 - accuracy: 0.6185 - val\_loss: 0.8721 - val\_accuracy: 0.7425  
Epoch 2/10  
704/704 [=====] - 105s 150ms/step - loss: 0.7927 - accuracy: 0.7543 - val\_loss: 0.5985 - val\_accuracy: 0.8063  
Epoch 3/10  
704/704 [=====] - 107s 152ms/step - loss: 0.6607 - accuracy: 0.7871 - val\_loss: 0.5356 - val\_accuracy: 0.8193  
Epoch 4/10  
704/704 [=====] - 105s 150ms/step - loss: 0.5994 - accuracy: 0.8028 - val\_loss: 0.4987 - val\_accuracy: 0.8321  
Epoch 5/10  
704/704 [=====] - 107s 152ms/step - loss: 0.5549 - accuracy: 0.8135 - val\_loss: 0.4756 - val\_accuracy: 0.8380  
Epoch 6/10

```
704/704 [=====] - 105s 149ms/step - loss: 0.5003 - accurac
y: 0.8285 - val_loss: 0.4470 - val_accuracy: 0.8462
Epoch 8/10
704/704 [=====] - 105s 149ms/step - loss: 0.4789 - accurac
y: 0.8331 - val_loss: 0.4308 - val_accuracy: 0.8495
Epoch 9/10
704/704 [=====] - 104s 148ms/step - loss: 0.4611 - accurac
y: 0.8390 - val_loss: 0.4276 - val_accuracy: 0.8513
Epoch 10/10
704/704 [=====] - 105s 149ms/step - loss: 0.4436 - accurac
y: 0.8428 - val_loss: 0.4179 - val_accuracy: 0.8547
625/625 [=====] - 9s 14ms/step - loss: 0.4295 - accuracy:
0.8498
625/625 [=====] - 9s 14ms/step
Learning rate: 0.0001
Units: 512
Accuracy: 0.8498
Precision: 0.7788123184105319
Recall: 0.6996241869554957
Runtime: 1066.905986070633
```

```
Epoch 1/10
 23/704 [.....] - ETA: 1:43 - loss: 5.1569 - accuracy: 0.02
45
```

```
IOPub message rate exceeded.
The notebook server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub_msg_rate_limit`.
```

```
Current values:
NotebookApp.iopub_msg_rate_limit=1000.0 (msgs/sec)
NotebookApp.rate_limit_window=3.0 (secs)
```

704/704 [=====] - 108s 153ms/step - loss: 1.1737 - accuracy: 0.6679 - val\_loss: 0.8864 - val\_accuracy: 0.7388  
Epoch 6/10  
704/704 [=====] - 108s 153ms/step - loss: 1.0775 - accuracy: 0.6886 - val\_loss: 0.8173 - val\_accuracy: 0.7539  
Epoch 7/10  
704/704 [=====] - 108s 153ms/step - loss: 1.0016 - accuracy: 0.7065 - val\_loss: 0.7691 - val\_accuracy: 0.7651  
Epoch 8/10  
704/704 [=====] - 108s 154ms/step - loss: 0.9438 - accuracy: 0.7199 - val\_loss: 0.7270 - val\_accuracy: 0.7756  
Epoch 9/10  
704/704 [=====] - 110s 156ms/step - loss: 0.8899 - accuracy: 0.7318 - val\_loss: 0.6949 - val\_accuracy: 0.7844  
Epoch 10/10  
704/704 [=====] - 108s 153ms/step - loss: 0.8533 - accuracy: 0.7414 - val\_loss: 0.6677 - val\_accuracy: 0.7923  
625/625 [=====] - 9s 15ms/step - loss: 0.6708 - accuracy: 0.7898  
625/625 [=====] - 9s 14ms/step  
Learning rate: 1e-05  
Units: 512  
Accuracy: 0.78985  
Precision: 0.7265378257011776  
Recall: 0.5978188654641512  
Runtime: 1097.3781945705414

Learning rate: 0.01  
Units: 128  
Accuracy: 0.8477  
Precision: 0.8043123489373756  
Recall: 0.6988400193387927  
Runtime: 919.8214240074158

Learning rate: 0.001  
Units: 128  
Accuracy: 0.85255  
Precision: 0.7867328400171182  
Recall: 0.707327582008939  
Runtime: 915.5359442234039

Learning rate: 0.0001  
Units: 128  
Accuracy: 0.843  
Precision: 0.786125959157098  
Recall: 0.6869673416696267  
Runtime: 901.746609210968

Learning rate: 1e-05  
Units: 128  
Accuracy: 0.73465  
Precision: 0.7162887549989162  
Recall: 0.47558125347175484  
Runtime: 885.5641613006592

Learning rate: 0.01  
Units: 256  
Accuracy: 0.85055  
Precision: 0.8134135377605697  
Recall: 0.7079299811099503  
Runtime: 990.635219335556

Learning rate: 0.001  
Units: 256

Accuracy: 0.85725  
Precision: 0.7893143274662201  
Recall: 0.7183029748170656  
Runtime: 909.7056584358215

Learning rate: 0.0001  
Units: 256  
Accuracy: 0.8464  
Precision: 0.7865026115391711  
Recall: 0.6944006943451094  
Runtime: 911.2518379688263

Learning rate: 1e-05  
Units: 256  
Accuracy: 0.76695  
Precision: 0.6844942561186504  
Recall: 0.5549426168871804  
Runtime: 991.9178223609924

Learning rate: 0.01  
Units: 384  
Accuracy: 0.84665  
Precision: 0.7905247884057616  
Recall: 0.7130944499335667  
Runtime: 991.4726612567902

Learning rate: 0.001  
Units: 384  
Accuracy: 0.8541  
Precision: 0.7530003563189905  
Recall: 0.7165125146157162  
Runtime: 984.3839774131775

Learning rate: 0.0001  
Units: 384  
Accuracy: 0.8433  
Precision: 0.7891048209621495  
Recall: 0.7014753587907294  
Runtime: 1054.4654290676117

Learning rate: 1e-05  
Units: 384  
Accuracy: 0.77725  
Precision: 0.690018898958171  
Recall: 0.5786048612733493  
Runtime: 1114.2524206638336

Learning rate: 0.01  
Units: 512  
Accuracy: 0.84525  
Precision: 0.7942784310891914  
Recall: 0.7128416009952518  
Runtime: 1098.0348672866821

Learning rate: 0.001  
Units: 512  
Accuracy: 0.85215  
Precision: 0.7637378673224909  
Recall: 0.7207188433737077  
Runtime: 1113.2304723262787

Learning rate: 0.0001  
Units: 512  
Accuracy: 0.8498



Precision: 0.7788123184105319  
Recall: 0.6996241869554957  
Runtime: 1066.905986070633

Learning rate: 1e-05  
Units: 512  
Accuracy: 0.78985  
Precision: 0.7265378257011776  
Recall: 0.5978188654641512  
Runtime: 1097.3781945705414

# Graph process

In [1]:

```
# Import packages
import numpy as np
import matplotlib.pyplot as plt
```

In [2]:

```
## SVM RESULTS

# Linear
l_c = [0.0010, 100.00]

l_acc = [0.7264, 0.7264]
l_pre = [0.5819, 0.5819]
l_rec = [0.5627, 0.5627]
l_tim = [16.308, 16.308] # in mins

# non-linear
c = [0.0010, 0.0178, 0.3162, 5.6234, 100.00]

# rbf
r_acc = [0.0547, 0.5330, 0.7315, 0.7877, 0.7774]
r_pre = [0.0078, 0.2271, 0.5956, 0.6843, 0.6745]
r_rec = [0.0162, 0.2063, 0.4716, 0.5996, 0.6010]
r_tim = [83.476, 371.14, 32.134, 25.899, 30.315] # in mins

# poly
p_acc = [0.0551, 0.2600, 0.6809, 0.7869, 0.7811]
p_pre = [0.0474, 0.3639, 0.5755, 0.6781, 0.6726]
p_rec = [0.0175, 0.1091, 0.3982, 0.5922, 0.6062]
p_tim = [148.38, 64.932, 45.097, 30.960, 29.848] # in mins
```

In [3]:

```
## SVM GRAPHS

# Produce Line Plots

fig, axs = plt.subplots(2, 2, figsize=(15, 7.5))

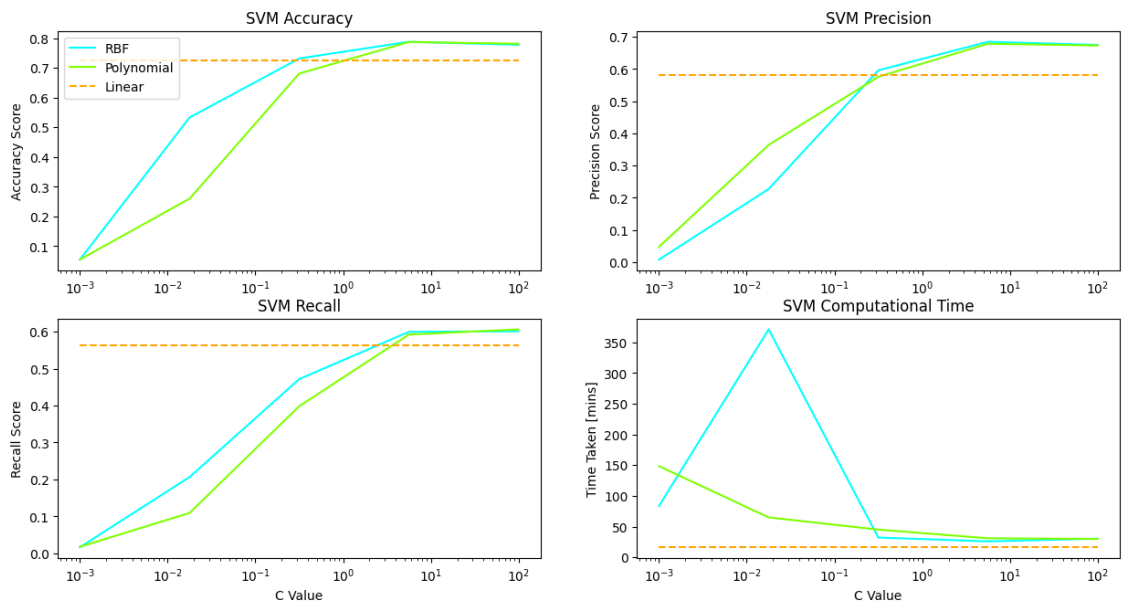
# ACCURACY
axs[0, 0].plot(c, r_acc, color='cyan')
axs[0, 0].plot(c, p_acc, color='chartreuse')
axs[0, 0].plot(l_c, l_acc, color='orange', linestyle='dashed')
axs[0, 0].legend(('RBF', 'Polynomial', 'Linear'), loc='upper left')
axs[0, 0].set_title('SVM Accuracy')
#axs[0, 0].set_xlabel('C Value')
axs[0, 0].set_ylabel('Accuracy Score')
axs[0, 0].set_xscale('log')

# PRECISION
axs[0, 1].plot(c, r_pre, color='cyan')
axs[0, 1].plot(c, p_pre, color='chartreuse')
axs[0, 1].plot(l_c, l_pre, color='orange', linestyle='dashed')
#axs[0, 1].legend(('RBF', 'Polynomial', 'Linear'), loc='upper left')
axs[0, 1].set_title('SVM Precision')
#axs[0, 1].set_xlabel('C Value')
axs[0, 1].set_ylabel('Precision Score')
axs[0, 1].set_xscale('log')

# RECALL
axs[1, 0].plot(c, r_rec, color='cyan')
axs[1, 0].plot(c, p_rec, color='chartreuse')
axs[1, 0].plot(l_c, l_rec, color='orange', linestyle='dashed')
#axs[1, 0].legend(('RBF', 'Polynomial', 'Linear'), loc='upper left')
axs[1, 0].set_title('SVM Recall')
axs[1, 0].set_xlabel('C Value')
axs[1, 0].set_ylabel('Recall Score')
axs[1, 0].set_xscale('log')

# COMPUTATIONAL TIME
axs[1, 1].plot(c, r_tim, color='cyan')
axs[1, 1].plot(c, p_tim, color='chartreuse')
axs[1, 1].plot(l_c, l_tim, color='orange', linestyle='dashed')
#axs[1, 1].legend(('RBF', 'Polynomial', 'Linear'), loc='upper left')
axs[1, 1].set_title('SVM Computational Time')
axs[1, 1].set_xlabel('C Value')
axs[1, 1].set_ylabel('Time Taken [mins]')
axs[1, 1].set_xscale('log')

#plt.savefig('svm_experiment_singlelegend.pdf', bbox_inches='tight', pad_inches=0.3, dpi = 500)
```



In [4]:

## ## CNN RESULTS

```
units = [128.00, 256.00, 384.00, 512.00]
```

### # LR 1e-2

```
e2acc = [0.8477, 0.8506, 0.8467, 0.8453]
e2pre = [0.8043, 0.8134, 0.7905, 0.7943]
e2rec = [0.6988, 0.7079, 0.7131, 0.7128]
e2tim = [15.330, 16.511, 16.525, 18.301] # in mins
```

### # LR 1e-3

```
e3acc = [0.8526, 0.8573, 0.8541, 0.8522]
e3pre = [0.7867, 0.7893, 0.7530, 0.7637]
e3rec = [0.7073, 0.7183, 0.7165, 0.7207]
e3tim = [15.259, 15.162, 16.406, 18.554] # in mins
```

### # LR 1e-4

```
e4acc = [0.8430, 0.8464, 0.8433, 0.8498]
e4pre = [0.7861, 0.7865, 0.7891, 0.7788]
e4rec = [0.6870, 0.6944, 0.7015, 0.6996]
e4tim = [15.029, 15.188, 17.574, 17.782] # in mins
```

### # LR 1e-5

```
e5acc = [0.7347, 0.7670, 0.7773, 0.7899]
e5pre = [0.7163, 0.6845, 0.6900, 0.7265]
e5rec = [0.4756, 0.5549, 0.5786, 0.5978]
e5tim = [14.759, 16.532, 18.571, 18.290] # in mins
```

In [5]:

```
## CNN GRAPHS

# Produce Line Plots

fig, axs = plt.subplots(2, 2, figsize=(15, 7.5))

# ACCURACY
axs[0, 0].plot(units, e2acc, color='cyan')
axs[0, 0].plot(units, e3acc, color='chartreuse')
axs[0, 0].plot(units, e4acc, color='orange')
axs[0, 0].plot(units, e5acc, color='violet')
axs[0, 0].legend(('Learning rate: 1e-02', 'Learning rate: 1e-03', 'Learning rate: 1e-04', 'Learning rate: 1e-05'))
axs[0, 0].set_title('CNN Accuracy')
#axs[0, 0].set_xlabel('Number of Dense Units')
axs[0, 0].set_ylabel('Accuracy Score')
#axs[0, 0].set_xscale('log')

# PRECISION
axs[0, 1].plot(units, e2pre, color='cyan')
axs[0, 1].plot(units, e3pre, color='chartreuse')
axs[0, 1].plot(units, e4pre, color='orange')
axs[0, 1].plot(units, e5pre, color='violet')
#axs[0, 1].legend(('Learning rate: 1e-02', 'Learning rate: 1e-03', 'Learning rate: 1e-04', 'Learning rate: 1e-05'))
axs[0, 1].set_title('CNN Precision')
#axs[0, 1].set_xlabel('Number of Dense Units')
axs[0, 1].set_ylabel('Precision Score')
#axs[0, 1].set_xscale('log')

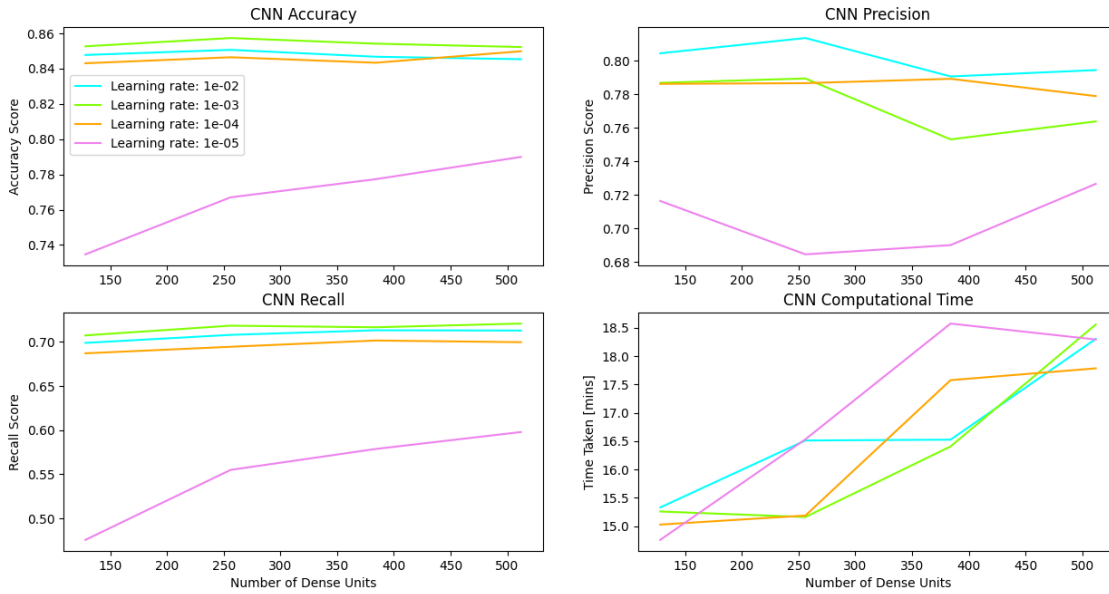
# RECALL
axs[1, 0].plot(units, e2rec, color='cyan')
axs[1, 0].plot(units, e3rec, color='chartreuse')
axs[1, 0].plot(units, e4rec, color='orange')
axs[1, 0].plot(units, e5rec, color='violet')
#axs[1, 0].legend(('Learning rate: 1e-02', 'Learning rate: 1e-03', 'Learning rate: 1e-04', 'Learning rate: 1e-05'))
axs[1, 0].set_title('CNN Recall')
axs[1, 0].set_xlabel('Number of Dense Units')
axs[1, 0].set_ylabel('Recall Score')
#axs[1, 0].set_xscale('log')

# COMPUTATIONAL TIME
axs[1, 1].plot(units, e2tim, color='cyan')
axs[1, 1].plot(units, e3tim, color='chartreuse')
axs[1, 1].plot(units, e4tim, color='orange')
axs[1, 1].plot(units, e5tim, color='violet')
#axs[1, 1].legend(('Learning rate: 1e-02', 'Learning rate: 1e-03', 'Learning rate: 1e-04', 'Learning rate: 1e-05'))
axs[1, 1].set_title('CNN Computational Time')
axs[1, 1].set_xlabel('Number of Dense Units')
axs[1, 1].set_ylabel('Time Taken [mins]')
#axs[1, 0].set_xscale('log')

plt.savefig('cnn_experiment_singlelegend.pdf', bbox_inches='tight', pad_inches=0.3, dpi = 500)
```

Out[5]:

```
Text(0, 0.5, 'Time Taken [mins]')
```



## ## MLP RESULTS

```
n_neu = [50.000, 66.000, 83.000, 100.00, 116.00, 133.00, 150.00, 166.00, 183.00, 200.00]
```

### # RELU

```
r_acc = [0.7868, 0.7922, 0.8030, 0.8025, 0.8086, 0.8125, 0.8125, 0.8147, 0.8154, 0.8126]
r_pre = [0.6636, 0.6652, 0.6812, 0.6863, 0.6939, 0.6982, 0.7013, 0.6930, 0.6940, 0.7021]
r_rec = [0.6185, 0.6344, 0.6429, 0.6486, 0.6528, 0.6641, 0.6544, 0.6728, 0.6648, 0.6722]
r_tim = [0.2138, 0.2233, 0.2327, 0.2923, 0.3003, 0.4704, 0.5358, 0.4253, 0.4250, 0.4445] # in mins
```

### # SIGMOID

```
s_acc = [0.7568, 0.7700, 0.7840, 0.7901, 0.7984, 0.8029, 0.8053, 0.8106, 0.8093, 0.8116]
s_pre = [0.6096, 0.6250, 0.6600, 0.6758, 0.6876, 0.6874, 0.6887, 0.7030, 0.6879, 0.6993]
s_rec = [0.5539, 0.5751, 0.5995, 0.6182, 0.6297, 0.6400, 0.6351, 0.6487, 0.6430, 0.6541]
s_tim = [0.2808, 0.2752, 0.2767, 0.3148, 0.3440, 0.3292, 0.3756, 0.3317, 0.3293, 0.5226] # in mins
```

### # TANH

```
t_acc = [0.7815, 0.7937, 0.7999, 0.8033, 0.8076, 0.8150, 0.8069, 0.8122, 0.8155, 0.8084]
t_pre = [0.6459, 0.6639, 0.6921, 0.6831, 0.6873, 0.7001, 0.6904, 0.6900, 0.7004, 0.7131]
t_rec = [0.6050, 0.6251, 0.6362, 0.6559, 0.6567, 0.6618, 0.6547, 0.6677, 0.6664, 0.6643]
t_tim = [0.2080, 0.2269, 0.2690, 0.2566, 0.2838, 0.2688, 0.2873, 0.2582, 0.4052, 0.4425] # in mins
```

In [7]:

```
## MLP GRAPHS

# Produce Line Plots

fig, axs = plt.subplots(2, 2, figsize=(15, 7.5))

# ACCURACY
axs[0, 0].plot(n_neu, r_acc, color='cyan')
axs[0, 0].plot(n_neu, s_acc, color='chartreuse')
axs[0, 0].plot(n_neu, t_acc, color='orange')
axs[0, 0].legend(('ReLU', 'Sigmoid', 'Tanh'), loc='upper left')
axs[0, 0].set_title('MLP Accuracy')
#axs[0, 0].set_xlabel('Number of Hidden Neurons')
axs[0, 0].set_ylabel('Accuracy Score')

# PRECISION
axs[0, 1].plot(n_neu, r_pre, color='cyan')
axs[0, 1].plot(n_neu, s_pre, color='chartreuse')
axs[0, 1].plot(n_neu, t_pre, color='orange')
#axs[0, 1].legend(('ReLU', 'Sigmoid', 'Tanh'), loc='upper left')
axs[0, 1].set_title('MLP Precision')
#axs[0, 1].set_xlabel('Number of Hidden Neurons')
axs[0, 1].set_ylabel('Precision Score')





# RECALL
axs[1, 0].plot(n_neu, r_rec, color='cyan')
axs[1, 0].plot(n_neu, s_rec, color='chartreuse')
axs[1, 0].plot(n_neu, t_rec, color='orange')
#axs[1, 0].legend(('ReLU', 'Sigmoid', 'Tanh'), loc='upper left')
axs[1, 0].set_title('MLP Recall')
axs[1, 0].set_xlabel('Number of Hidden Neurons')
axs[1, 0].set_ylabel('Recall Score')

# COMPUTATIONAL TIME
axs[1, 1].plot(n_neu, r_tim, color='cyan')
axs[1, 1].plot(n_neu, s_tim, color='chartreuse')
axs[1, 1].plot(n_neu, t_tim, color='orange')
#axs[1, 1].legend(('ReLU', 'Sigmoid', 'Tanh'), loc='upper left')
axs[1, 1].set_title('MLP Computational Time')
axs[1, 1].set_xlabel('Number of Hidden Neurons')
axs[1, 1].set_ylabel('Time Taken [mins]')

#plt.savefig('mlp_experiment_singlelegend.pdf', bbox_inches='tight', pad_inches=0.3, dpi = 500)
```

Out[7]:

```
Text(0, 0.5, 'Time Taken [mins]')
```



# Overall best models

#	acc.	pre.	rec.	tim.	#
svm	0.8042	0.7042	0.6340	0.7657	
cnh	0.8522	0.7637	0.7207	0.6185	
mlp	0.8126	0.7021	0.6722	0.0148	

# time is % of 30 mins



In [9]:

```
measures = ("Accuracy Score", "Precision Score", "Recall Score", "Time Taken [30mins]")
scores = {
    'Support Vector Machine': (0.804, 0.704, 0.634, 0.766),
    'Convolutional Neura': (0.852, 0.764, 0.721, 0.619),
    'MLP': (0.813, 0.702, 0.672, 0.015),
}

x = np.arange(len(measures)) # the label locations
width = 0.18 # the width of the bars
multiplier = 0
colors = ['cyan', 'chartreuse', 'orange']

fig, ax = plt.subplots(figsize=(12, 6))

for attribute, measurement in scores.items():
    offset = width * multiplier
    rects = ax.bar(x + offset, measurement, width, label=attribute, color=colors[multiplier])
    ax.bar_label(rects, padding=3)
    multiplier += 1

# Add some text for labels, title and custom x-axis tick labels, etc.
ax.set_title('Comparison of Evaluation Metrics for All Models')
ax.set_xticks(x + width, measures)
ax.legend(loc='upper left', ncols=3)
ax.set_ylim(0, 1)

plt.savefig('best_models.pdf', bbox_inches='tight', pad_inches=0.3, dpi = 500) #save plot as pdf
```

Out[9]:

(0.0, 1.0)

